

Bayesian neural network estimation of next-to-leading-order cross sections

by

René Alexander Ask

THESIS

for the degree of

MASTER OF SCIENCE



Faculty of Mathematics and Natural Sciences
University of Oslo

Spring 2022

Bayesian neural network estimation of next-to-leading-order cross sections

René Alexander Ask

© 2022 René Alexander Ask

Bayesian neural network estimation of next-to-leading-order cross sections

<http://www.duo.uio.no/>

Printed: Reprosentralen, University of Oslo

Abstract

This is my abstract.

Acknowledgments

Acknowledgments yo

Contents

Introduction	1
1 The Physics Problem	3
1.1 Computation of Beyond the Standard Model Cross Sections	3
1.2 Tmp	3
2 Bayesian Formulation of Machine Learning	5
2.1 The Core of Machine Learning	5
2.1.1 Loss Functions	5
2.1.2 Regularization	6
2.1.3 Optimization	6
2.2 Bayes' theorem	6
2.3 Bayesian Framework for Machine Learning	7
2.4 Bayesian Inference	8
3 Markov Chain Monte Carlo	11
3.1 Expectation Values and the Typical Set	11
3.1.1 The Typical Set	11
3.1.2 The Target Density and Bayesian Applications	12
3.2 Markov Chains and Markov Transitions	12
3.2.1 Ideal Markov Chains	13
3.2.2 Pathologies	13
3.2.3 Geometric Ergodicity and Convergence Diagnostics	13
3.3 Metropolis-Hastings	13
3.3.1 The Proposal Distribution	14
3.4 Gibbs Sampling	15
4 Hamiltonian Monte Carlo	17
4.1 Hamiltonian Dynamics	17
4.1.1 Leapfrog integration	18
4.2 Generating a Proposal State	19
4.3 The Potential Energy Function in Bayesian Machine Learning Applications	21
4.4 Limitations of Hamiltonian Monte Carlo	21
5 Adaptive Hamiltonian Monte Carlo	23
5.1 The No-U-Turn Sampler	23
5.1.1 Generation of States and the Stopping Criterion	24
5.1.2 Selecting Candidate Points	25
5.1.3 Efficiently Implementing the No-U-Turn Sampler	25
5.1.4 Computational Cost of The No-U-Turn Sampler	27

5.2	Adapting the Step Size	27
5.2.1	Adapting a General Parameter Using a Dual Averaging Scheme	28
5.2.2	Setting the Step Size in Hamiltonian Monte Carlo	28
6	Bayesian Neural Networks	29
6.1	Neural Networks	29
6.1.1	Basic Mathematical Structure	29
6.1.2	Backpropagation	30
6.1.3	Regularization in Neural Networks	31
6.2	Activation Functions	31
6.2.1	Sigmoid and Tanh	31
6.2.2	ReLU	32
6.2.3	Swish	32
6.3	Bayesian learning of Neural Networks using Monte Carlo Samplers	32
6.3.1	What <i>is</i> Bayesian learning of Neural Networks?	32
6.3.2	The Potential Energy Function of Neural Networks	33
6.3.3	Practical Training of Bayesian Neural Networks	33
6.3.4	Training Algorithm of Bayesian Neural Networks	34
7	Numerical Experiments	35
7.1	The Dataset	35
7.1.1	Data Generation	35
7.1.2	Data Scaling and Transformations	35
7.2	Performance Metrics	35
7.2.1	Relative Error	35
7.2.2	Standardized Residuals	36
7.3	Results	36
7.3.1	Benchmarks of Hyperparameters	36
7.3.1.1	Baseline Model	36
7.3.1.2	Pretraining	36
7.3.1.3	Burn-in length	36
7.3.1.4	Number of model parameters	36
7.3.2	Neutralino-Neutralino Cross Sections	36
8	Discussion	39
	Conclusion	42
	Appendices	43
	Appendix A	45
A.1	Appendix 1 title	45

List of Algorithms

3.1	Metropolis-Hastings	14
3.2	Gibbs sampling	15
4.1	Leapfrog Integration	18
4.2	Vectorized Leapfrog Integration	19
4.3	Hamiltonian Monte Carlo	21
5.1	Helper Function Used with the NUTS Sampler	26
5.2	The NUTS Sampler	27
6.1	Backpropagation: Forward pass	30
6.2	Backpropagation: Backward pass	31

Introduction

Motivation, context and problem.

- Contrast variational inference to Bayesian inference with MCMC methods.

Outline of the Thesis

Give outline of thesis

Chapter 1

The Physics Problem

In this chapter, we shall discuss the need for application of machine learning models in physics problems and why this is a necessary solution.

1.1 Computation of Beyond the Standard Model Cross Sections

The Standard Model of particle physics (SM) is a successful fundamental theory that describes fundamental particles and their interactions. Despite its success, however, it has a few limitations on its own which has led physicists to propose extensions to the model to explain physics that SM cannot. One such family of extensions is called *supersymmetry*.

1.2 Tmp

- n_i : Målte events (kollisjoner) som oppfyller et sett (kalt signalregion) med kriterier (“cuts”).
- b_i : Bakgrunnen. Estimert SM bidrag for samme signal region.
- s_i : BSM estimert bidrag for signalregionen med et sett med parameterverdier for en ny BSM modell (i.e SUSY). Den er regnet ut ved

$$s_i = \sigma \epsilon_i A_i \mathcal{L}, \quad (1.1)$$

der σ er tverrsnittet som måler sannsynligheten for at en “ny” prosess skjer, ϵ_i er detektor effektivitet, A_i er akseptans og \mathcal{L} er integrert luminositet over data brukt i søket.

- Statistisk analyse gjøres med å regne ut Poisson likelihood

$$\mathcal{L}(s, b, n) = \frac{e^{-(s+b)} (s+b)^n}{n!}. \quad (1.2)$$

og en test statistikk

$$q = -2 \ln \frac{\mathcal{L}(s, b, n)}{\mathcal{L}(s=0, b, n)}. \quad (1.3)$$

Chapter 2

Bayesian Formulation of Machine Learning

In this chapter we will introduce the notion of *Bayesian machine learning* (Bayesian ML). We will start from the classical view of ML and reformulate it in terms of Bayesian concepts. We will only concern ourselves with so-called supervised ML models used to solve supervised regression tasks as it is the only class of problems of interest in this thesis. We will first introduce the core of ML and its constituent ingredients. From this we transition to Bayes' theorem and a Bayesian framework for ML. Finally we discuss Bayesian inference.

2.1 The Core of Machine Learning

The basic conceptual framework of a supervised machine learning problem is as follows. Assume a dataset D is a sequence of N datapoints $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$, where $x^{(i)} \in \mathbb{R}^p$ is the set of *features* and $y^{(i)} \in \mathbb{R}^d$ is the *target*. The next ingredient is to assume the targets can be decomposed as

$$y = f(x) + \epsilon, \quad (2.1)$$

for some true function $f : \mathbb{R}^p \rightarrow \mathbb{R}^d$ (also known as the *ground truth*), where $\epsilon \in \mathbb{R}^d$ is introduced to account for random noise. The objective is to learn $f(x)$ from the dataset. To this end, we choose a *model class* $\hat{f}(x; \theta)$ parameterized by a model parameters $\theta \in \mathbb{R}^m$, combined with a procedure to infer an estimate of the parameters $\hat{\theta}$ such that the model is as close to $f(x)$ as possible. Formally, this means choosing a *metric* \mathcal{L} to quantify the error, called a *loss* function (or a *cost* function, but we will adopt the former term in line with the terminology used in the TensorFlow framework), and minimize it with respect to the parameters of the model to obtain $\hat{\theta}$ using an optimization algorithm. For brevity, we will denote the output of a model class as $\hat{y}^{(i)} \equiv \hat{f}(x^{(i)}; \theta)$.

2.1.1 Loss Functions

For regression problems, two loss functions \mathcal{L} are commonly chosen. The first is the *residual sum of squares* (RSS) given by

$$\mathcal{L}_{\text{RSS}} \equiv \text{RSS} = \sum_{i=1}^N \left\| y^{(i)} - \hat{y}^{(i)} \right\|_2^2, \quad (2.2)$$

where $\|\cdot\|_2$ denotes the L^2 -norm. The second is the *mean squared error* (MSE), defined as

$$\mathcal{L}_{\text{MSE}} \equiv \text{MSE} = \frac{1}{N} \sum_{i=1}^N \left\| y^{(i)} - \hat{y}^{(i)} \right\|_2^2. \quad (2.3)$$

For optimization purposes, they yield equivalent optimal parameters $\hat{\theta}$, at least in principle.

2.1.2 Regularization

With datasets of limited size, *overfitting* can pose a problem, yielding models that generalize poorly because they become overly specialized to the dataset on which $\hat{\theta}$ is inferred. The implication is that the predicted target on unseen data is unlikely to be correct. This occurs especially if the model is too complex. One strategy to overcome this, is to tack on a regularization term to the loss-function. By *regularization*, we mean an additional term that limits the size of the allowed parameter space. Hence, regularization imposes a constraint on the optimization problem.

The two most commonly used regularization terms are L^2 -regularization, which adds a term to the loss function as

$$\mathcal{L} = \mathcal{L}_0 + \frac{\lambda}{2} \|\theta\|_2^2, \quad (2.4)$$

where λ is the so-called *regularization strength*, which is what we call a *hyperparameter*, and \mathcal{L}_0 is a loss function with no regularization term. The second is L^1 -regularization, which yields a loss

$$\mathcal{L} = \mathcal{L}_0 + \frac{\lambda}{2} \|\theta\|_1. \quad (2.5)$$

The terms *penalize* large values of θ , effectively shrinking the allowed parameter space. The larger the value of the regularization strength λ , the smaller the allowed parameter space becomes.

More generally, we can decomposed our full loss function as

$$\mathcal{L}(x, y, \theta) = \mathcal{L}_0 + R(\lambda_1, \dots, \lambda_r, \theta), \quad (2.6)$$

where $R(\theta)$ is a linear combination of L^p -regularization terms where λ_i are the expansion coefficients which are all treated as hyperparameters. L^p -regularization terms is defined by the L^p -norm

$$\|x\|_p = (|x_1|^p + \dots + |x_m|^p)^{1/p}, \quad x \in \mathbb{R}^m. \quad (2.7)$$

In practice, we typically use a single form of L^p -regularization but nothing stops us from constructing complicated regularization terms in theory.

2.1.3 Optimization

Once a model class and loss function is chosen, an *optimizer* or *optimization algorithm* must be chosen. By this, we mean an algorithm that uses the loss function and the model class, and minimizes the loss with respect to the model parameters to yield an estimate of $\hat{\theta}$. Regardless of which optimization algorithm we employ, we seek

$$\hat{\theta} = \arg \min_{\theta} \mathcal{L}. \quad (2.8)$$

In this thesis, optimization plays a smaller role in the inference of model parameters than in classical ML because we do not seek a single estimate $\hat{\theta}$ in most Bayesian applications. We shall nevertheless utilize such algorithms for some parts but for another purpose. One of the most popular optimizers in the deep learning community is ADAM [1] which we will mainly use when optimization is needed.

2.2 Bayes' theorem

Our goal is to reformulate ML in terms of Bayesian concepts. The backbone of Bayesian ML is *Bayes' theorem* [2]. The theorem can be formulated as

$$p(\theta|D) = \frac{p(D|\theta)p(\theta)}{p(D)}, \quad (2.9)$$

where D is observed data and θ denotes the parameters of the model. Here $p(\theta)$ is called the *prior* distribution and embodies our prior knowledge of θ before any new observations are considered. $p(D|\theta)$ is called the *likelihood* function and provides the relative probability of observing D for a fixed value of θ . It need not be normalized to unity, which is why it only provides relative “probabilities”. The *posterior* distribution $p(\theta|D)$ models our belief about θ after the data D is observed. Finally, $p(D)$ is called the *evidence* which we may regard as the normalization constant of the posterior such that posterior integrates to unity over parameter space. In the context of Bayesian ML, the evidence will not be an interesting quantity as it will not turn up as part of any algorithms. Moreover, it is typically intractable for sufficiently large parameter spaces. It is therefore common to write Bayes’ theorem as

$$p(\theta|D) \propto p(D|\theta)p(\theta), \quad (2.10)$$

which we too shall adopt.

2.3 Bayesian Framework for Machine Learning

The Bayesian framework for ML differs somewhat in approach to its classical counterpart. We define a model class in the same way as before. Choosing a loss function is substituted with choosing a likelihood function and a prior. Minimization of the loss function is replaced with maximization of the likelihood function or the posterior distribution. In fact, The Bayesian framework introduces several ways to infer an estimate for the optimal model parameters [3].

1. *Maximum Likelihood Estimation* (MLE): The optimal parameters $\hat{\theta}$ are inferred by

$$\hat{\theta} = \arg \max_{\theta} p(D|\theta), \quad (2.11)$$

meaning we choose $\hat{\theta}$ as the mode of the likelihood function. This is equivalent to maximizing the log-likelihood (since log is a monotonic function), i.e.

$$\hat{\theta} = \arg \max_{\theta} \log p(D|\theta). \quad (2.12)$$

2. *Maximum-A-Posteriori* (MAP): This estimate of $\hat{\theta}$ is defined as

$$\hat{\theta} = \arg \max_{\theta} p(\theta|D), \quad (2.13)$$

meaning we choose $\hat{\theta}$ as a mode of the posterior distribution.

3. *Bayes’ estimate*: The estimate of $\hat{\theta}$ is chosen as the expectation of the posterior,

$$\hat{\theta} = \mathbb{E}_{p(\theta|D)}[\theta] = \int d\theta \theta p(\theta|D). \quad (2.14)$$

The connection between classical and Bayesian ML can be understood from what follows. First, let us assume that each datapoint $(x^{(i)}, y^{(i)})$ is identically and independently distributed (i.i.d.). The likelihood function can then generally be written as

$$P(D|\theta) = \prod_{i=1}^N P(y^{(i)}|x^{(i)}, \theta). \quad (2.15)$$

For regression tasks, the standard choice of likelihood function is the *Gaussian*

$$p(y|x, \theta) = \exp \left(-\frac{1}{2\sigma^2} \|y - \hat{f}(x; \theta)\|_2^2 \right), \quad (2.16)$$

where σ is some hyperparameter typically chosen to be the same for every datapoint (x, y) . For the full dataset, we get

$$p(D|\theta) = \prod_{i=1}^N \exp \left(-\frac{1}{2\sigma^2} \left\| y^{(i)} - \hat{f}(x^{(i)}; \theta) \right\|_2^2 \right). \quad (2.17)$$

Now, consider the definition of MLE from eq. (2.11). It instructs us to maximize the expression in eq. (2.17). If we rewrite the likelihood function a bit

$$p(D|\theta) = \exp \left(-\frac{1}{2\sigma^2} \sum_{i=1}^N \left\| y^{(i)} - \hat{f}(x^{(i)}; \theta) \right\|_2^2 \right), \quad (2.18)$$

we can observe that maximization of the likelihood function simply amounts to minimization of the RSS and hence of the MSE, as can be seen by comparison with the expressions in eq. (2.2) and eq. (2.3).

We can go even further, by considering the MAP estimate. Let us introduce a Gaussian prior on the parameters such that

$$p(\theta) \propto \exp \left(-\frac{\lambda}{2} \|\theta\|_2^2 \right). \quad (2.19)$$

The posterior obtained from Bayes' theorem in eq. (2.10) by combining the prior introduced in eq. (2.19) and the likelihood function in eq. (2.17) is

$$p(\theta|D) \propto p(D|\theta)p(\theta) \propto \prod_{i=1}^N \exp \left(-\frac{1}{2\sigma^2} \left\| y^{(i)} - \hat{f}(x^{(i)}; \theta) \right\|_2^2 \right) \exp \left(-\frac{\lambda}{2} \|\theta\|_2^2 \right), \quad (2.20)$$

which we can rewrite as

$$p(\theta|D) \propto \exp \left(-\left[\frac{1}{2\sigma^2} \sum_{i=1}^N \left\| y^{(i)} - \hat{f}(x^{(i)}; \theta) \right\|_2^2 + \frac{\lambda}{2} \|\theta\|_2^2 \right] \right). \quad (2.21)$$

Maximization of this expression is equivalent to minimization of RSS or MSE with a L^2 -regularization term tacked on which can be seen by comparison with eq. (2.4). Obviously, we are missing a factor $1/N$ in front of the likelihood term which can be thought of as baked into the σ parameter. The natural generalization is that the posterior can be expressed as

$$p(\theta|D) \propto \exp(-\mathcal{L}), \quad (2.22)$$

for any loss function as in eq. (2.6). For a purpose that comes much later when we discuss Hamiltonian Monte Carlo, we can invert eq. (2.22)

$$\mathcal{L} = -\log Z - \log p(D|\theta) - \log p(\theta), \quad (2.23)$$

for some appropriate normalization constant Z . Assuming that the dataset consists of observations that are i.i.d, we get

$$\mathcal{L} = -\log Z - \sum_{i=1}^N p(y^{(i)}|x^{(i)}, \theta) - \log p(\theta). \quad (2.24)$$

Equation (2.24) will play an important role later on.

2.4 Bayesian Inference

We have seen that there is a straight forward connection between the Bayesian framework and the classical view of ML by looking at estimators $\hat{\theta}$. In regression tasks, however, we are seldom interested

in a single estimate of the model parameter. Instead we seek to obtain the posterior distribution from which we can infer other quantities. In applications where the model class is sufficiently complex, direct computation of the posterior is not feasible. Instead, we must settle with an approximate posterior distribution which we construct using Monte Carlo Markov chains (MCMC) methods. The discussion of such methods is allocated to chapter 3. For now we assume that there exists a way to generate samples $\theta \sim p(\theta|D)$. We approximate the posterior by sampling a set of model parameters $\{\theta^{(1)}, \dots, \theta^{(n)}\}$ where $\theta^{(t)} \sim p(\theta|D)$, yielding an *empirical* posterior distribution.

We will primarily use the posterior to compute two classes of mathematical objects. The first is the *predictive distribution* of a target y^* given an input x^* . The predictive distribution can be expressed as

$$p(y^*|x^*, D) = \int d\theta p(y^*|x^*, \theta)p(\theta|D). \quad (2.25)$$

Equation (2.25) is generally intractable since we cannot exactly compute the posterior. The predictive distribution is therefore approximated by generating a set of predictions using the empirical posterior distribution. That is, we indirectly sample from $p(y^*|x^*, D)$ by computation of $\hat{f}(x^*; \theta^{(t)})$ for $t = 1, \dots, n$. In other words, the empirical predictive distribution is generated as follows.

$$\begin{aligned} \theta^{(t)} &\sim p(\theta|D), \\ f(x^*; \theta^{(t)}) &\sim p(y^*|x^*, \theta). \end{aligned} \quad (2.26)$$

The second class is expectation values with respect to the posterior distribution, which for a target function $f(\theta)$ is defined as

$$\mathbb{E}_{p(\theta|D)}[f] = \int d\theta f(\theta)p(\theta|D). \quad (2.27)$$

An important example of eq. (2.27) is the expectation value of the predictive distribution, which will be the expectation of the model class with respect to the posterior

$$\hat{y} \equiv \mathbb{E}_{p(\theta|D)}[\hat{f}(x; \theta)] = \int d\theta \hat{f}(x; \theta)p(\theta|D). \quad (2.28)$$

Equation (2.27) must be approximated since we cannot hope to evaluate the posterior $p(\theta|D)$. Even if we could, we will be working with sufficiently large parameters spaces such that the integral itself is intractable in any case. Approximation of expectation values is done using MCMC methods which is the subject of the next chapter.

Chapter 3

Markov Chain Monte Carlo

In this chapter, we will discuss fundamental ideas pertaining to *Markov Chain Monte Carlo* (MCMC) methods. We shall confine the discussion to continuous sample spaces which is the kind needed in this thesis. We will commence with a discussion of expectation values and an important notion called the *typical set*. We will then define and discuss Markov chains and Markov transitions after which we shall discuss Metropolis-Hastings sampling and its limitations. Finally we will look at Gibbs sampling. We will adopt a geometric view where possible to provide a natural transition to Hamiltonian Monte Carlo and the No-U-Turn sampler in the two following chapters.

3.1 Expectation Values and the Typical Set

Consider a *target probability density* $\pi(\theta)$ and an m -dimensional sample space Θ where $\theta \in \Theta$. Consider $f(\theta)$ to be an arbitrary smooth function of θ . The *expectation value* of $f(\theta)$ with respect to the density $\pi(\theta)$ is then defined as

$$\mathbb{E}_{\pi}[f] = \int d\theta \pi(\theta) f(\theta). \quad (3.1)$$

We shall interchangeably refer to expectation values simply as *expectations*. We will call the function f we seek to compute the expectation of as the *target function*. For all but a few simple densities, evaluation of eq. (3.1) is impossible analytically. To complicate things further, numerical evaluation with numerical integration techniques of the expectation in high-dimensional spaces quickly becomes computationally infeasible as the dimensionality increases, due to limited computational resources. Even worse, we may not even be able to write down the expression of $\pi(\theta)$ explicitly. Fortunately, it is unlikely that the entire sample space contribute significantly to the expectation. If we could somehow pick out the points in sample space that *does* contribute, only knowing $\pi(\theta)$ up to a normalization constant, we could make approximate computations of expectations tractable.

For most purposes, we are interested in the expectation of more than a single target function. For example, in Bayesian applications, we are often interested in both the mean and variance of a quantity which introduces the need for several target functions. Thus the numerical method should not depend on the target function in question. Instead the focus should be laid on the contribution from $\pi(\theta)d\theta$ to the integrand. The objective of MCMC methods is to efficiently sample points from regions of sample space where this quantity is non-negligible. This region of sample space is called the *typical set* [4].

3.1.1 The Typical Set

For simplicity, we can divide a sample space into three regions with respect to the target density $\pi(\theta)$.

1. High-probability density region. These are regions in the neighborhood of a mode of the target density. In general, as the dimensionality increases, the contribution from $\pi(\theta)d\theta$ becomes negligible here unless the volume in the region is significant enough.
2. The typical set. This refers to the regions in which $\pi(\theta)d\theta$ provides a non-negligible contribution to any expectation. This may be thought of as the high-probability region of the sample space since $\pi(\theta)d\theta$ is proportional to probability of a volume $d\theta$ in the neighborhood of θ .
3. Low-probability density regions. These are regions far away from any mode of the density. This region, too, will generally yield negligible contributions to the integrand even if the volume is large.

Although the notion of a typical set can be formalized precisely, we will intentionally operate with this somewhat imprecise definition. For our purposes, it suffices to use it merely as a conceptual notion to evaluate the quality of the samples generated by an MCMC chain.

3.1.2 The Target Density and Bayesian Applications

In the chapter on Bayesian ML, we mentioned that we could not compute the evidence term of Bayes' theorem in realistic applications and thus were only concerned with a proportionality relationship $p(\theta|D) \propto p(D|\theta)p(\theta)$. Thus any MCMC methods we are interested in cannot require that the $\pi(\theta)$ is normalized to unity. We only require that the density is smooth and that

$$0 < \int d\theta \pi(\theta) < \infty. \quad (3.2)$$

Sometimes we may refer to the target density as the *target distribution*. In Bayesian applications, we assume that $\pi(\theta) = p(D|\theta)p(\theta)$ such that $p(\theta|D) \propto \pi(\theta)$.

3.2 Markov Chains and Markov Transitions

Since direct evaluation of eq. (3.1) in most applications is intractable, we seek to approximately evaluate it by generating samples $\theta^{(t)}$ from the typical set using *Markov chains*. A Markov chain is a sequence of points $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(n)}$ generated sequentially using a random map called a *Markov transition*. A Markov transition is a conditional probability density $T(\theta'|\theta)$ that yields the probability of transition from a point θ to θ' . The Markov transition is also called a *Markov kernel* which is a special case of a *transition kernel*. The latter is the term we will adopt because it is the term used by TensorFlow Probability.

An arbitrary transition kernel is not useful because the generated Markov chain is unlikely to have any relation to the target distribution of interest. To generate a useful Markov chain, we must use a transition kernel that preserves the target distribution. The condition that ensures this is

$$\pi(\theta) = \int d\theta' \pi(\theta') T(\theta|\theta'). \quad (3.3)$$

The condition is formally called *detailed balance*. The interpretation of the condition is that the Markov chain is reversible.

We can start from any θ and use the transition kernel to produce a set of new states. The distribution generated by the Markov chain should be distributed according the target distribution regardless of which point we used to generate the chain from, given a long enough chain. A more important fact is that as long as this condition is satisfied, the Markov chain will converge to and stay within the typical set.

The standard approach to approximate eq. (3.1) is then with the MCMC *estimator*

$$\hat{f}_N = \frac{1}{N} \sum_{t=1}^N f(\theta^{(t)}). \quad (3.4)$$

For large enough N , the estimator can be shown to converge to the true expectation such that $\lim_{N \rightarrow \infty} \hat{f}_N = \mathbb{E}_\pi[f]$. Obviously, the knowledge that the estimator will asymptotically converge to the true expectations is of limited use when restricted to a practical computation in which only a finite chain can be generated. We must therefore understand the properties of finite Markov chains so we can efficiently use them to approximate eq. (3.1).

3.2.1 Ideal Markov Chains

In order to understand the behaviour of finite Markov chains, we should first consider the behaviour of ideal Markov chains. An ideal Markov chain can be divided into three phases.

1. A convergence phase. The Markov chain is initiated from some point θ and the initially generated sequence lies in a region outside the typical set. Estimators evaluated using this part of the sequence are highly biased, meaning inclusion of these points will lead to an estimator that lies relatively far away from the true expectation.
2. An exploration phase. The Markov chain has reached the typical set and begins its first “traversal” of it. In this phase, estimators will rapidly converge towards the true expectations.
3. A saturation phase. At this point, the Markov chain has explored most of the typical set and convergence of the estimators slow down significantly.

The ideal evaluation of estimators thus only use the parts of the Markov chain generated in the second and third phase, discarding the the chain generated in the first phase. The notion of discarding the chain from the first phase is called *burn-in* or *mixing*. To most efficiently approximate eq. (3.1), we should really only use points generated in the exploration phase. Using points from the saturation phase does not hurt the estimators but yield diminishing returns with respect to computational resources.

3.2.2 Pathologies

Unfortunately, many target distributions embody typical sets with pathological regions where *any* transition kernel that obey eq. (3.3) is not sufficient to *efficiently* explore the typical set. Geometrically, this can be regions in the typical set in which the target distribution rapidly changes. The pathological regions can be completely ignored by the chain for much of the exploration, leading to poor convergence and thus biased estimators. However, as long as the transition kernel satisfies detailed balance, we know for a fact that the estimators *must* converge eventually. Consequentially, the Markov chain will be stuck near pathological regions for long periods to compensate before it rapidly explores other parts of the typical set. This behaviour can be repeated, which makes estimators oscillate. Regardless of when the MCMC chain is terminated, the estimator will likely be biased due to this oscillating behaviour.

3.2.3 Geometric Ergodicity and Convergence Diagnostics

Generation of ideal Markov chains is guaranteed if the transition kernel satisfies *geometric ergodicity* [5], a *Central Limit Theorem* for the MCMC estimators. However, in most cases it is impossible to check that the condition is satisfied. Instead one uses a statistical quantity known as the *potential scale reduction factor* \hat{R} [6]. The ideal value is $\hat{R} = 1$. For values far away from this target, it is unlikely that geometric ergodicity is satisfied. The Rule-of-thumb is to assume convergence if $\hat{R} < 1.1$ [7].

3.3 Metropolis-Hastings

Construction of a transition kernel that ensures convergence to the typical set of the target distribution is a non-trivial problem in general. Fortunately, the Metropolis-Hastings algorithm provides a general

solution that lets us construct a transition kernel with this property [8, 9]. The algorithm consist of two components; a proposal of a new state and a correction step called the *Metropolis correction*. Given a state θ , we propose a new state θ' by adding a random perturbation to the initial state. The correction step rejects a proposed state that moves away from the typical set of the target distribution and accepts proposals that stay within it. The proposed state is formally sampled from a *proposal distribution* $q(\theta'|\theta)$. The probability of accepting the proposed state given the initial state, fittingly called the *acceptance probability*, is

$$a(\theta'|\theta) = \min \left(1, \frac{q(\theta|\theta')\pi(\theta')}{q(\theta'|\theta)\pi(\theta)} \right). \quad (3.5)$$

A particularly neat feature of the acceptance probability in eq. (3.5) is that it can be calculated in Bayesian applications because the evidence term cancels out. These steps are summarized in algorithm 3.1.

Algorithm 3.1 Metropolis-Hastings

```

function MetropolisHastings( $\theta$ )
  Sample  $\theta' \sim q(\theta'|\theta)$ 
   $a(\theta'|\theta) \leftarrow \min \left( 1, \frac{q(\theta|\theta')\pi(\theta')}{q(\theta'|\theta)\pi(\theta)} \right)$ 
  Sample  $u \sim \text{Uniform}(0, 1)$ .
  if  $a(\theta'|\theta) \geq u$  then
     $\theta \leftarrow \theta'$  ▷ Accept transition
  else
     $\theta \leftarrow \theta$  ▷ Reject transition
  end if
  return  $\theta$ 
end function

```

3.3.1 The Proposal Distribution

There are many valid choices of proposal distributions. A common choice is a Gaussian distribution $q(\theta'|\theta) = \mathcal{N}(\theta'|\theta, \Sigma)$, where Σ is the *covariance matrix* of the normal distribution used to generate the perturbation of the initial state. This is typically chosen to be the identity matrix $\Sigma = I$.

We will refer to the Metropolis-Hastings algorithm with this proposal distribution as *random walk Metropolis*. More precisely, this means that a proposed state is given by

$$\theta' = \theta + \delta, \quad (3.6)$$

where $\delta \sim \mathcal{N}(0, \Sigma)$. This distribution is symmetric such that $q(\theta'|\theta) = q(\theta|\theta')$, implying that the acceptance probability reduces to

$$a(\theta'|\theta) = \min \left(1, \frac{\pi(\theta')}{\pi(\theta)} \right). \quad (3.7)$$

Hence, evaluation of the acceptance probability only require that we evaluate the target distribution at the initial state and the proposed state.

The random walk Metropolis algorithm does suffer from slow convergence to, and exploration of, the typical set in high-dimensional spaces. This can be understood because of the following. As we increase the dimension of the sample space, the volume outside of the typical set becomes increasingly larger than the volume of the typical set itself. This implies with increasing probability that a random

perturbation of an arbitrary initial state will cause the proposed state to lie outside the typical set for a fixed covariance matrix. We can compensate for this flaw by reducing the values of Σ_{ij} , but this will slow down exploration of the sample space. The slow exploration also leads to a Markov chain where consecutive samples embody a relatively large measure of correlation. The quality of the resulting Markov chain tarnishes and successive samples must be discarded in order to properly evaluate eq. (3.1). This process of discarding correlated successive samples in a Markov chain is called *thinning*. Fortunately, there exists a solution; *gradient-informed* exploration of the sample space, manifested in the form of *Hamiltonian Monte Carlo*. This algorithm is a special case of a Metropolis-Hastings algorithm in which the proposal distribution $q(\theta'|\theta)$ is a special one utilizing Hamiltonian dynamics and Gibbs sampling to produce a new proposal state θ' . This is the topic of the next chapter.

3.4 Gibbs Sampling

The final standard MCMC algorithm we need is the *Gibbs* sampler. It plays a small part of the sampling in HMC and so we should therefore briefly discuss it. It is a MCMC sampling method used for multi-variate probability densities, and so is only meaningful to discuss for $d > 1$ dimensions. Suppose $\theta^{(t)}$ represents the parameters at iteration t . The next sample $\theta^{(t+1)}$ in the Markov chain is drawn according to some chosen conditional distribution P depending on the previous and current sample as follows

$$\theta_i^{(t+1)} \sim P(\theta_i | \theta_1^{(t+1)}, \dots, \theta_{i-1}^{(t+1)}, \theta_{i+1}^{(t)}, \dots, \theta_m^{(t)}). \quad (3.8)$$

We may summarize this as a function in algorithm 3.2 which given an initial state $\theta^{(t)}$ returns a new state $\theta^{(t+1)}$ sampled according to eq. (3.8).

Algorithm 3.2 Gibbs sampling

```

function Gibbs( $\theta^{(t)}$ )
  for  $i = 1, \dots, d$  do
    Sample  $\theta_i^{(t+1)} \sim P(\theta_i | \theta_1^{(t+1)}, \dots, \theta_{i-1}^{(t+1)}, \theta_{i+1}^{(t)}, \dots, \theta_m^{(t)})$ .
  end for
  return  $\theta^{(t+1)} = (\theta_1^{(t+1)}, \dots, \theta_m^{(t+1)})$ .
end function

```

Chapter 4

Hamiltonian Monte Carlo

In this chapter, we will explore the details of Hamiltonian Monte Carlo. It is a Markov chain Monte Carlo method that uses gradient-informed steps to generate a proposal state for Metropolis correction. This is achieved by usage of Hamiltonian dynamics which allow gradient-informed exploration by treating the model parameters as “coordinates” of a fictitious physical system, and introducing auxiliary variables representing its momenta. The coordinates and momenta are required to obey a particular set of coupled differential equations called Hamilton’s equations. The differential equations cannot in general be solved exactly and are instead simulated. The particular kind of numerical method used to achieve this is called the Leapfrog integrator. At the end of a simulation, a new set of coordinates and momenta will be generated, which is regarded as the proposal state to undergo Metropolis correction. If accepted, we keep the proposed coordinates as the next parameter in the Markov chain. Otherwise, the initial coordinates assume this role. The auxiliary momenta is discarded and resampled on each iteration as they play no important role for the actual Markov chain.

We begin by presenting Hamiltonian dynamics and the Leapfrog integrator. Once established we show how the framework is used to construct an MCMC method. Next, we will see how to apply the method to Bayesian machine learning models before we finalize the chapter with a discussion on some limitations of the method.

4.1 Hamiltonian Dynamics

Hamiltonian dynamics [10] is a formulation of classical mechanics that allows us to compute the time evolution of a physical system. The fundamental mathematical object of the theory is the *Hamiltonian* H which governs the time evolution of the *coordinates* $q = (q_1, \dots, q_d)$ and *momenta* $p = (p_1, \dots, p_d)$ of the system. The $2d$ -dimensional space defined by the points (q, p) is called *phase-space*. The precise relationship is formulated by *Hamilton’s equations*

$$\frac{dq_i}{dt} = \frac{\partial H}{\partial p_i}, \quad \frac{dp_i}{dt} = -\frac{\partial H}{\partial q_i}, \quad \text{for } i = 1, \dots, d. \quad (4.1)$$

The objective is to use the $2d$ coupled differential equations in eq. (4.1) to find $(q(t), p(t))$ given some initial condition $(q(0), p(0))$ where t represents time. A system governed by Hamilton’s equations is called a *Hamiltonian system*. For the purpose of constructing a MCMC method, we need not consider the most general theory of Hamiltonian dynamics and we will therefore refrain from doing so. We shall confine our focus to Hamiltonians which can be decomposed as

$$H(q, p) = V(q) + K(p), \quad (4.2)$$

where V is the *potential energy* and K is the *kinetic energy* of the system. The particular kind of Hamiltonian in eq. (4.2) corresponds to the total energy of the system. A key feature is that this

Hamiltonian is *conserved* through time. This observation follows from

$$\frac{dH}{dt} = \sum_i \left(\frac{dq_i}{dt} \frac{\partial H}{\partial q_i} + \frac{dp_i}{dt} \frac{\partial H}{\partial p_i} \right) = \sum_i \left(\frac{\partial H}{\partial p_i} \frac{\partial H}{\partial q_i} - \frac{\partial H}{\partial q_i} \frac{\partial H}{\partial p_i} \right) = 0. \quad (4.3)$$

Thus any solution $(q(t), p(t))$ will be confined to a hyperplane defined by the Hamiltonian and the initial condition.

Evolving a Hamiltonian system from some initial point $(q(0), p(0))$ is in general a non-trivial task. Exact solutions can only be computed for simple systems. To arm ourselves with a robust MCMC method, then, we must employ a numerical method to approximate the solutions. To this end, there is a class of numerical methods called *symplectic integrators* that take advantage of the underlying geometry enforced by Hamilton's equations which allow accurate solutions over long time periods at a lower computational cost than typical higher-order methods such as fourth-order Runge-Kutta. The particular symplectic integrator used in HMC is called the *Leapfrog integrator* which we shall discuss next.

4.1.1 Leapfrog integration

The Leapfrog integrator [11] is used in HMC to integrate eq. (4.1) to generate new proposal states. First assume that we *discretize* the time-coordinate t into discrete time coordinates defined by an initial time t_0 and a *step size* ϵ which defines the distance between each time coordinate. The k -th time coordinate can be generated by

$$t_k = t_0 + k\epsilon. \quad (4.4)$$

To please mathematicians, we introduce functions \hat{q} and \hat{p} to represent the discretized approximations to the exact solution $(q(t), p(t))$. From an initial point $(q(t_0), p(t_0))$, we simulate the system to obtain approximate values of the exact solution at discrete times t_1, \dots, t_n .

Consider a single Leapfrog step from a point $(\hat{q}(t), \hat{p}(t))$. Its approximation to $(q(t + \epsilon), p(t + \epsilon))$ is then computed as formulated in algorithm 4.1.

Algorithm 4.1 Leapfrog Integration

```

function Leapfrog( $V, q, p, \epsilon$ )
  for  $i = 1, \dots, d$  do
     $p'_i \leftarrow p_i - \frac{\epsilon}{2} \frac{\partial V(q)}{\partial q_i}$ 
     $q'_i \leftarrow q_i + \frac{\epsilon}{m_i} p'_i$ 
     $p'_i \leftarrow p'_i - \frac{\epsilon}{2} \frac{\partial V(q')}{\partial q_i}$ 
  end for
  return  $(q', p')$ 
end function

```

Note the introduction of the *masses* m_i . For now they may simply be regarded as some constants belonging to the Hamiltonian system. When used in HMC, it is common to set all masses $m_i = 1$ from which we can formulate the algorithm in vectorized form, as seen in algorithm 4.2.

Algorithm 4.2 Vectorized Leapfrog Integration

```

function VectorizedLeapfrog( $V, q, p, \epsilon$ )
   $p' \leftarrow p - \frac{\epsilon}{2} \nabla_q V(q)$ 
   $q' \leftarrow q + \epsilon p'$ 
   $p' \leftarrow p' - \frac{\epsilon}{2} \nabla_q V(q')$ 
  return  $(q', p')$ 
end function

```

4.2 Generating a Proposal State

Our next objective is to understand how we connect an arbitrary target distribution $\pi(\theta)$ to Hamiltonian dynamics. In this section we will weave these together and show how we generate a new proposal state θ' which will undergo a Metropolis correction. The results discussed in this section can be understood as representing the proposal distribution $q(\theta'|\theta)$ used during the Metropolis-Hastings step, as we summarized in algorithm 3.1.

The fundamental assumption we make is that the target distribution can be expressed in terms of a canonical distribution over coordinate space

$$\pi(q) \propto \exp\{-V(q)\}, \quad (4.5)$$

where q represents the model parameters θ . We will stick to this convention to avoid confusion and utilize the formulation of Hamiltonian dynamics discussed hitherto. Once we want to apply it in a Bayesian ML context, we simply replace $q \rightarrow \theta$. From eq. (4.5), we can find the potential energy function in terms of the target distribution

$$V(q) = -\log \pi(q), \quad (4.6)$$

up to a constant. Hence once the target distribution is known, we use eq. (4.6) to obtain the potential energy of the system.

We now turn to the problem of constructing the Hamiltonian so we can utilize Hamilton's equations. To achieve this, we must introduce *auxilliary* momenta p so we can define a kinetic energy function and evolve the system through what we may regard as *fictitious time* t . The momenta are sampled from some distribution of our own choice. We can proceed in the same way as we did with the potential energy function and express the momentum distribution in terms of a canonical distribution over momentum space

$$\pi(p) \propto \exp\{-K(p)\}, \quad (4.7)$$

such that

$$K(p) = -\log \pi(p), \quad (4.8)$$

up to a constant. The commonly chosen expression for kinetic energy is the one found in classical physics

$$K(p) = \sum_{i=1}^d \frac{p_i^2}{2m_i}, \quad (4.9)$$

from which the canonical distribution is inferred to be

$$\pi(p) \propto \exp\left\{-\sum_{i=1}^d \frac{p_i^2}{2m_i}\right\} = \prod_{i=1}^d \exp\left\{-\frac{p_i^2}{2m_i}\right\}. \quad (4.10)$$

Hence, with the kinetic energy from eq. (4.9), we sample each momentum independently from a Gaussian distribution with zero mean and variance $\sigma_i^2 = m_i$.

Now that we understand how we specify the potential energy for a given target distribution and the kinetic energy of the auxilliary momenta, we can formulate the full canonical distribution over phase-space as

$$\pi(q, p) = \pi(q)\pi(p) \propto \exp\{-V(q)\} \exp\{-K(p)\} = \exp\{-H(q, p)\}. \quad (4.11)$$

We are naturally just interested in generating a new coordinate q' . Using Hamilton's equations with the Hamiltonian implied by eq. (4.11), we can simulate the fictitious Hamiltonian system using Hamilton's equation in eq. (4.1) to generate a new state (q', p') . The proposal state is then obtained by the projection map $(q', p') \mapsto q'$.

As stated in the beginning of this section, we may regard the details outlined here as an elaborate explanation of the proposal distribution $q(\theta'|\theta)$. The final keypoint to consider is how we can make it symmetric so that we only need to evaluate $\pi(q', p')$ at the Metropolis step using eq. (3.7). It can be shown that we only need two additional steps. We must randomly choose to sample forwards or backwards in time. The second step is the negate the momenta at the end of the generation of the state, $p \mapsto -p$. The acceptance probability can then be computed as

$$a = \min\left(1, \frac{\pi(q', p')}{\pi(q, p)}\right) = \min(1, \exp\{-[H(q', p') - H(q, p)]\}). \quad (4.12)$$

But this should always be evaluated to $a = 1$ if H is indeed conserved. But the catch is that the dynamics is only approximated using the Leapfrog integrator. The best the integrator can do is conserve H on average, with its value oscillating about the initial value.

Before we summarize the algorithm in a neat manner, we shall briefly outline it conceptually.

1. Given an initial state q , we randomly sample the auxilliary momenta p from the distribution in eq. (4.10) to generate an initial condition (q, p) to use with Hamilton's equations.
2. We randomly choose to simulate the system forwards or backwards in time by sampling a variable $v \sim \text{Uniform}(\{-1, 1\})$ from which the step size is set as $v\epsilon$. Forwards in time is represented by $v = 1$ and backwards in time is represented by $v = -1$.
3. Perform L Leapfrog steps using algorithm 4.1 for a total trajectory length of ϵL to produce a proposal point (q', p') .
4. Perform a Metropolis-Hastings correction on the proposal state to accept or reject it.
5. Project the phase-space point onto coordinate space and return q' if accepted, or q if rejected, in the previous step.

This essentially summarizes the practical steps of HMC. The introduction of randomly simulating forwards and backwards in time is to ensure that the algorithm is reversible and obeys the detailed balance condition discussed in chapter 3. To please mathematicians once more, we must really reverse the sign of the final momenta as well, but since we shall use a Gaussian distribution, changing the sign of the momenta makes no difference to the value of the kinetic energy. To generate a Markov chain by this procedure, we simply feed the returned coordinate state back in to the machinery and reiterate. The HMC scheme is summarized in algorithm 4.3.

Algorithm 4.3 Hamiltonian Monte Carlo

```

function HMCstep( $q, H, L, \epsilon$ )
  Sample  $p \sim \mathcal{N}(0, \text{diag}(m_1, \dots, m_d))$  ▷ Sample auxilliary momenta
  Sample  $v \sim \text{Uniform}(\{-1, 1\})$ . ▷ Randomly choose direction in time.
   $(q', p') \leftarrow (q, p)$  ▷ Initialize the initial state.
  for  $l = 1, \dots, L$  do ▷ Simulate Hamiltonian system for  $L$  Leapfrog steps.
     $(q', p') \leftarrow \text{Leapfrog}(q', p', v\epsilon)$ 
  end for
   $a = \min(1, \exp\{-[H(q', p') - H(q, p)]\})$  ▷ Compute acceptance probability
  Sample  $u \sim U(0, 1)$  ▷ Uniform distribution on  $(0, 1)$ .
  if  $a \geq u$  then ▷ Perform Metropolis-Hastings correction
     $q \leftarrow q'$  ▷ Accept proposed state.
  else
     $q \leftarrow q$  ▷ Reject proposed state.
  end if
  return  $q$ 
end function

```

4.3 The Potential Energy Function in Bayesian Machine Learning Applications

We seek to use HMC in a Bayesian ML application. It is therefore important to discuss a general way to construct the potential energy function in such applications. First, recall from chapter 2 in eq. (2.22) that the posterior could in general be written as

$$p(\theta|D) \propto \exp\{-\mathcal{L}(\theta)\}, \quad (4.13)$$

where \mathcal{L} was some loss function in the classical ML sense. However, we do not need the evidence term and simply sample from the target distribution $\pi(\theta) = p(D|\theta)p(\theta)$ instead. Comparison with eq. (4.5) makes it clear that the potential energy function simply is \mathcal{L} . Combining this with eq. (2.24), lets us conclude that the general expression for the potential energy is

$$\mathcal{L} = -\log p(D|\theta) - \log p(\theta), \quad (4.14)$$

up to a constant. If we assume all N datapoints are i.i.d. we can recast it as eq. (2.24), that is

$$\mathcal{L} = -\sum_{i=1}^N \log p(y^{(i)}|x^{(i)}, \theta) - \log p(\theta). \quad (4.15)$$

4.4 Limitations of Hamiltonian Monte Carlo

Although HMC is effective at exploring the state space we wish to sample from, it suffers from the need to hand-tune the trajectory length ϵL . Poor choices of ϵ and L can lead to poor results. On one hand, if the trajectory length is too short, exploration of the state space will be limited which makes HMC behave like a random-walk. Suppose we fix the trajectory length to a finite, but sufficiently large value. If the step size ϵ is too large, it can lead to instabilities in the leapfrog integrator, while if its chosen to be too small, it will perform far too many iterations to make the algorithm worthwhile. Tuning these parameters requires preliminary runs for the problem at hand and analysis of so-called *trace statistics*, which essentially measures the quality of the generated Markov chain.

In the next chapter, we will look at algorithms that adaptively sets the trajectory length of HMC, namely the No-U-Turn sampler combined with dual-averaging of the step size, which allows us to overcome these limitations and more effectively sample from the target distribution without the need for hand-tuning and analysis of trace statistics, or reliance on heuristics.

Chapter 5

Adaptive Hamiltonian Monte Carlo

Hamiltonian Monte Carlo is considered a state-of-the-art sampler that produces successive samples with low correlation that may lie far apart in sample space, but suffers the need for manual tuning of the trajectory length ϵL . In this chapter, we will explore improvements that adaptively adjust the trajectory length. This is achieved by means of adapting both the number of Leapfrog steps L using a improved sampler called the *No-U-Turn* (NUTS) sampler, and an adaptive scheme for setting the step size ϵ using a *dual averaging* algorithm. We will closely follow the treatment in the original paper [12] but adapt the notation to fit with the previous chapter.

We will start off with a discussion on how to adapt the number of Leapfrog steps using NUTS. At a high-level, NUTS starts from an initial state (q, p) and simulates the Hamiltonian dynamics of the system. This is done in the following way. Leapfrog steps are performed either forwards or backwards in time, first with a single Leapfrog step, then two Leapfrog steps, then followed by four Leapfrog steps and so on. This reiteration of the simulation is performed until the the path traced out starts to double back towards the initial point. The states traced out can be regarded as a *balanced binary tree* where each node represents a phase-space state produced by the Leapfrog integrator during the simulation. The proposal state is sampled at random from these nodes to be fed through a Metropolis correction step.

We will finalize the chapter with the dual averaging scheme for adaptively setting the step size using with the Leapfrog integrator. The algorithm is a modified version of a dual averaging scheme presented by Nesterov in [13].

5.1 The No-U-Turn Sampler

The No-U-Turn sampler augments standard HMC by introduction of a *slice variable* u which is sampled according to

$$u \sim p(q, p) = \text{Uniform}(u; [0, \exp\{-H(q, p)\}]) . \quad (5.1)$$

This implies the conditional distribution for (q, p) given the slice variable as

$$p(q, p|u) = \text{Uniform}\left(q, p; \left\{q', p' \mid \exp\{-H(q', p')\} \geq u\right\}\right) . \quad (5.2)$$

Consequentially, we have a joint distribution

$$p(q, p, u) \propto \mathbb{I}[u \in [0, \exp\{-H(q, p)\}]] , \quad (5.3)$$

where $\mathbb{I}[\cdot]$ evaluates to 1 if its argument is true and 0 otherwise. Integrating with respect to u yields the marginal distribution over phase-space

$$p(q, p) = \int p(q, p, u) du \propto \int_0^{\exp\{-H(q, p)\}} du = \exp\{-H(q, p)\} . \quad (5.4)$$

which is the target distribution we use in standard HMC.

Let \mathcal{B} be the set of all states traced out by the Leapfrog integrator used in HMC. Let $\mathcal{C} \subseteq \mathcal{B}$ be the *candidate set* of all candidate states (q, p) from \mathcal{B} that obey detailed balance. The candidate set is deterministically constructed from \mathcal{B} by introducing a conditional distribution $p(\mathcal{B}, \mathcal{C} | q, p, u, \epsilon)$ with the following conditions imposed:

1. All elements of \mathcal{C} are volume perserving. This effectively translates to $p((q, p) | (q, p) \in \mathcal{C}) \propto p(q, p)$.
2. The current state must be included in \mathcal{C} , i.e $p((q, p) \in \mathcal{C} | q, p, u, \epsilon) = 1$.
3. Any state $(q', p') \in \mathcal{C}$, must be in the slice defined by u . Mathmatically, this is expressed as

$$p\left(u \leq \exp\{-H(q, p)\} \middle| (q', p') \in \mathcal{C}\right) = 1.$$

4. If $(q, p) \in \mathcal{C}$ and $(q', p') \in \mathcal{C}$, then for any \mathcal{B} we impose $p(\mathcal{B}, \mathcal{C} | q, p, u, \epsilon) = p(\mathcal{B}, \mathcal{C} | q', p', u, \epsilon)$. Thus any point in \mathcal{C} is equally likely, expressing a uniform distribution over the candidate set.

5.1.1 Generation of States and the Stopping Criterion

Up until this point, we have not yet described precisely how the states in \mathcal{B} are generated, nor why and when to stop its generation. As we briefly described in the introduction to this chapter, NUTS computes trajectories in phase-space until the trajectory starts to double back on itself. First running one Leapfrog step, then two Leapfrog steps, then four Leapfrog steps and so on. Each such step is run either forwards or backwards in fictitious time, chosen at random. If the directions in forwards, it starts from the state at the front of the total generated trajectory. If the direction is backwards in time, it starts from the state corresponding to the tail of the trajectory. The successive state produced by the Leapfrog integrator at each such step is collected and stored. This generates a collection of states which we represent with \mathcal{B} .

We can regard the process as building a balanced binary tree where each node correspond to a state traced out by the Leapfrog integrator. The initial node is defined to represent the tree at height $j = 0$. Given a balanced binary tree at height j , we use the last endpoint from the last simulation to run 2^j Leapfrog steps in the direction of $v_j \sim \text{Uniform}(\{-1, 1\})$, where $v_j = 1$ represents forwards in time and $v_j = -1$ represents backwards in time. If $v_j = 1$, the old tree of height j becomes the left half of the new tree and the 2^j states traced out by the Leapfrog integrator becomes the right half of the new tree of height $j + 1$. If $v_j = -1$, the old tree becomes the right half and the new tree becomes the left half.

We cannot continue this generation process forever, of course. At some point we must stop building \mathcal{B} and select all candidates states that will collectively define \mathcal{C} . The stopping criterions employed by the algorithm are

1. **Too large simulation error.** The slice variable u that was introduced as an augmentation of the HMC model requires $u \leq \exp\{-H(q, p)\}$ at any point during the simulation. The Leapfrog integrator can introduce some numerical error which leads to a violation of this relation. The No-U-Turn sampler loosens this requirement a bit to avoid an inefficient algorithm. Instead, it halts the simulation if

$$H(q, p) + \log u < \Delta_{\max}, \quad (5.5)$$

for some tolerance Δ_{\max} . The authors of the original paper recommend this to be set to $\Delta_{\max} = 1000$. This is not a concern as states that violate the slice condition will not be included in \mathcal{C} regardless.

2. **The “No-U-Turn” criterion.** If at any point during the simulation, the trajectory starts to move towards points the integrator has already visited, the simulation is terminated. We can

measure this by considering an initial point q and a point q' computed through integration. The change in their squared distance between the two points with respect to time is then proportional to

$$\frac{d}{dt} \frac{\|q' - q\|_2^2}{2} = (q' - q)^T \frac{d}{dt} (q' - q) = (q' - q)^T p', \quad (5.6)$$

where q is regarded as a constant and $p' = dq'/dt$. When the tree at height j is built, NUTS considers its $2^j - 1$ subtrees using eq. (5.6) in the following way. Consider an arbitrary chosen subtree out of the total $2^j - 1$ subtrees. Let (q^-, p^-) represent the state in its leftmost node and (q^+, p^+) its rightmost node. If these states for any subtree satisfy

$$(q^+ - q^-)^T p^+ < 0 \quad \text{or} \quad (q^+ - q^-)^T p^- < 0, \quad (5.7)$$

we terminate the simulation. The criterion can be interpreted as if we continue the simulation either forwards or backwards in time an infinitesimal duration dt , we would reduce the distance between q^+ and q^- . The criterion adds an additional cost of $2^{j+1} - 2$ inner products on top of what is required by HMC (two inner products per subtree). However, this additional cost is negligible for sufficiently complex models and/or large datasets. Computation of gradients of the potential energy function will in most cases be the dominating computational cost per iteration.

5.1.2 Selecting Candidate Points

Now that we know how to generate the states of \mathcal{B} , we turn our attention to how we select the candidate states that builds up \mathcal{C} . We need not write down an explicit expression for $p(\mathcal{B}, \mathcal{C} | q, p, u, \epsilon)$ since we can select the points in a way that reflects the four conditions discussed earlier. The first condition is automatically satisfied because the Leapfrog integrator is volume preserving. The second condition is satisfied as long as we include the initial state as part of \mathcal{C} . Condition three is satisfied as long as we only include states (q', p') that satisfy $u \leq \exp\{-H(q', p')\}$, that is, we only include points that lie in the slice defined by u . The fourth condition required that $p(\mathcal{B}, \mathcal{C} | q, p, u, \epsilon) = p(\mathcal{B}, \mathcal{C} | q', p', u, \epsilon)$. For any initial state $(q', p') \in \mathcal{B}$, there is at most one sequence of directions $\{v_j\}_{j=1}^J$ that can generate \mathcal{B} . Any state that cannot be used to recreate \mathcal{B} must be excluded from \mathcal{C} as these would violate detailed balance. This condition will be satisfied as long as any state that satisfy the stopping criterions in either eq. (5.5) or eq. (5.6) is excluded from \mathcal{C} . There are two cases which must be considered

1. Equation (5.5) was satisfied by a state or eq. (5.6) was satisfied by a subtree during the final doubling step. In this case, any element of \mathcal{B} that was added during the final doubling must be excluded from \mathcal{C} . This is because using any such state as an initial state to build \mathcal{B} is impossible because one of the stopping criterions will be met before one can fully rebuild the tree.
2. The doubling procedure is stopped because eq. (5.6) is satisfied by the leftmost and rightmost nodes of the full tree. In this case, no exclusion is necessary because from any such node in the tree, we can find a unique sequence of directions in time for any state in the tree from which we can recreate the entire tree before the stopping criterion is met.

So now we know how to select the candidate set from which we sample the final proposal state.

5.1.3 Efficiently Implementing the No-U-Turn Sampler

The details discussed so far imply that we must store 2^j position and momentum states for a tree of height j in memory. For sufficiently complex models or deep enough trees, the required memory footprint might be infeasible. We can overcome this limitation by noting that for any subtree $\mathcal{C} \subseteq \mathcal{C}$, we can express the uniform probability over \mathcal{C} as Then

$$p(q, p | \mathcal{C}) = \frac{1}{|\mathcal{C}|} = \frac{|\mathcal{C}_{\text{subtree}}|}{|\mathcal{C}|} \frac{1}{|\mathcal{C}_{\text{subtree}}|} = p((q, p) \in \mathcal{C}_{\text{subtree}} | \mathcal{C}) P(q, p | (q, p) \in \mathcal{C}_{\text{subtree}}, \mathcal{C}), \quad (5.8)$$

meaning it can be expressed as the product of the probability of selecting a node that belongs to the subtree times the probability of sampling (q, p) randomly from the states in that subtree. This observation can be practically applied to reduce the memory footprint as follows. Consider a candidate set \mathcal{C} representing a tree of height J . Any subtree of height $j > 0$ is built up of two smaller subtrees of height $j - 1$. For each smaller subtree, sample a pair (q, p) from $1/|\mathcal{C}_{\text{subtree}}|$ to represent each smaller subtree. We select one of these pairs and give it a weight proportional to how many elements of \mathcal{C} that reside in that subtree. This procedure can be performed repeated from $j = 1$ all the way up to the subtree that represent \mathcal{C} , meaning there is no need to store the entire tree corresponding to the candidate set. Since we select a single state (q, p) per subtree, the storage requirement goes as $\mathcal{O}(j)$ instead of $\mathcal{O}(2^j)$.

The algorithm is summarized by two components. A helper function **BuildTree** which is used to implicitly build the tree corresponding to the candidate set \mathcal{C} . It returns the leftmost node (q^-, p^-) and rightmost node (q^+, p^+) of the subtree it builds along with the state q' representing the subtree, with a weight n' and an indicator variable s' that tracks if a stopping criterion is met ($s' = 0$) or not ($s' = 1$). It is based on recursion to avoid explicitly storing this information and is readily easy to implement directly into a programming language such as Python. The second component is a one-step function similar to **HMCstep** in algorithm 4.3 which performs a single step with the NUTS sampler and produces a new sample q in the Markov chain it is used to generate. The function **NUTSstep** is listed in algorithm 5.2.

Algorithm 5.1 Helper Function Used with the NUTS Sampler

```

function BuildTree( $q, p, u, v, j, \epsilon$ )
  if  $j = 0$  then
     $q', p' \leftarrow \text{Leapfrog}(q, p, v\epsilon)$ 
     $n' \leftarrow \mathbb{I}[u \leq \exp\{-H(q', p')\}]$ 
     $s' \leftarrow \mathbb{I}[H(q', p') + \log u > \Delta_{\max}]$ 
    return  $q', p', q', p', q', n', s'$ 
  else
     $q^-, p^-, q^+, p^+, q', n', s' \leftarrow \text{BuildTree}(q, p, u, v, j - 1, \epsilon)$ 
    if  $s' = 1$  then
       $q^-, p^-, -, -, q'', n'', s'' \leftarrow \text{BuildTree}(q^-, p^-, u, v, j - 1, \epsilon)$ 
    else
       $-, -, q^+, p^+, q'', n'', s'' \leftarrow \text{BuildTree}(q^+, p^+, u, v, j - 1, \epsilon)$ 
    end if
     $q' \leftarrow q''$  with probability  $n''/(n' + n'')$ 
     $s' \leftarrow s'' \mathbb{I}[(q^+ - q^-)^T p^- \geq 0] \mathbb{I}[(q^+ - q^-)^T p^+ \geq 0]$ 
     $n' \leftarrow n' + n''$ 
    return  $q^-, p^-, q^+, p^+, q', n', s'$ 
  end if
end function

```

Algorithm 5.2 The NUTS Sampler

```

function NUTSstep( $q, H, \epsilon$ )
  Sample  $p \sim \mathcal{N}(0, I)$ .
  Initialize  $s = 1, q^\pm = q, p^\pm = p, j = 0, n = 1$ .
  Sample  $u \sim \text{Uniform}([0, \exp\{-H(q, p)\}])$ 
  while  $s = 1$  do
    Sample  $v_j \sim \text{Uniform}(\{-1, 1\})$ 
    if  $v_j = -1$  then
       $q^-, p^-, -, -, q', n', s' \leftarrow \text{BuildTree}(q^-, p^-, u, v, v_j, j, \epsilon)$ 
    else
       $-, -, q^+, p^+, q', n', s' \leftarrow \text{BuildTree}(q^+, p^+, u, v_j, j, \epsilon)$ 
    end if
    if  $s' = 1$  then
       $q \leftarrow q'$  with probability  $\min\{1, n'/n\}$ 
    end if
     $n \leftarrow n + n'$ 
     $s \leftarrow s' \mathbb{I}[(q^+ - q^-)^T p^- \geq 0] \mathbb{I}[(q^+ - q^-)^T p^+ \geq 0]$ 
     $j \leftarrow j + 1$ 
  end while
  return  $q$ 
end function

```

5.1.4 Computational Cost of The No-U-Turn Sampler

Let us discuss the computational cost of this algorithm. The algorithm demands $2^j - 1$ evaluations of $H(q, p)$ and its gradient. Moreover, an additional set of operations to determine if a stopping criterion is reached, which is of the order $\mathcal{O}(2^j)$. As argued earlier, though, the computational cost is comparable to standard HMC per leapfrog step when the model is sufficiently complex or the dataset large. However, in its current form it requires storage of 2^j positions and momenta, which for increasingly complex models, or deep balanced binary trees due to repeatedly calling `BuildTree` without any stopping criterion being met, may results in an intractibly large storage requirement. In the next section we shall explore a more efficient solution that reduces the memory footprint of the algorithm.

5.2 Adapting the Step Size

So far, we have seen how we can adapt the number of Leapfrog steps L used in HMC with the NUTS sampler but we have yet to explore how we can adapt the step size ϵ . After all, we sought to adapt the trajectory length ϵL itself. The adaptation scheme discussed in this section is a slightly modified version of the dual averaging scheme presented by Nesterov. The modifications, the specifics of which we will clarify later on, is better suited for MCMC problems while the methods presented in Nesterov's paper is designed for stochastic convex optimization problems.

5.2.1 Adapting a General Parameter Using a Dual Averaging Scheme

Assume H_t is a statistic that measures the some behaviour of an MCMC algorithm at iteration $t \geq 1$. Define its expectation with respect to a tunable parameter $x \in \mathbb{R}$ as

$$h(x) = \mathbb{E}_t[H_t|x] \equiv \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \mathbb{E}[H_t|x]. \quad (5.9)$$

Assuming that h is a nondecreasing function of x , the update

$$x_{t+1} \leftarrow x_t - \eta_t H_t, \quad (5.10)$$

ensures that $h(x_t) \rightarrow 0$ as long as the step size η_t satisfies the conditions

$$\sum_t \eta_t = \infty \quad \text{and} \quad \sum_t \eta_t^2 < \infty. \quad (5.11)$$

To satisfy the conditions in eq. (5.11), we can set the step size schedule as

$$\eta_t = t^{-k} \quad \text{for } k \in (1/2, 1]. \quad (5.12)$$

In practice, we want x_t to ensure that $h(x_t)$ is close to zero towards the end of the warm-up phase of the MCMC chain and freeze the value of x during the phase we gather samples from the chain. However, the optimal value of x may be widely different in the convergence phase and the exploration phase. We thus want our updates to give weight towards the end of the warm-up phase.

The objective is to find a value of x such that $h(x) = 0$. To this end, we apply the updates

$$\begin{aligned} x_{t+1} &\leftarrow \mu - \frac{\sqrt{t}}{\gamma} \frac{1}{t + t_0} \sum_{i=1}^t H_i, \\ \bar{x}_{t+1} &\leftarrow \eta_t x_{t+1} + (1 - \eta_t) \bar{x}_t, \end{aligned} \quad (5.13)$$

with $\bar{x}_1 = x_1$. Here μ is a free parameter that x_t is chosen to converge to, $\gamma > 0$ is a free parameter that controls the rate at which x converges to μ and $t_0 \geq 0$ is introduced to stabilize the initial iterations of the updates. The difference from Nestorov's scheme the introduction of t_0 .

5.2.2 Setting the Step Size in Hamiltonian Monte Carlo

In Hamiltonian Monte Carlo, we want to avoid a value for ϵ that is too small which would waste computational resources, or one that is too large which lead to low acceptance probabilities. We can avoid this by tuning ϵ to yield an average acceptance probability of some value chosen target δ . Under some strong conditions, it can be shown that the ideal value is $\delta \approx 0.65$ [14] but empirically any value $\delta \in (0.6, 0.9)$ works fine. To use the update rule in eq. (5.13) to tune ϵ with δ as target can be achieved by introduction of

$$H_t^{\text{HMC}} = \min \left\{ 1, \frac{\pi(q', p')}{\pi(q, p)} \right\}, \quad (5.14)$$

where (q', p') the proposed state and (q, p) is the initial point used by `HMCstep` as in algorithm 4.3. The observant reader will note that eq. (5.14) is simply the acceptance probability used during the Metropolis correction in `HMCstep` with eq. (4.11) as the target distribution. The updates are then performed with

$$H_t = \delta - H_t^{\text{HMC}} \quad \text{and} \quad x = \log \epsilon, \quad (5.15)$$

which will ensure that $h(\epsilon)$ converges to δ .

Chapter 6

Bayesian Neural Networks

6.1 Neural Networks

In this chapter, we will finally discuss the main topic of this thesis, *Bayesian neural networks* (BNNs). We will start off introducing the mathematical formalism of neural networks. We will then discuss the *backpropagation* algorithm, which is the standard algorithm used to compute the gradient of the model with respect to a specified loss. We will then finalize the chapter with how Bayesian learning of neural networks work. Fortunately, most of the groundwork is already laid, so we need only a mathematical description of the model and a Bayesian interpretation of it. We will stay general and assume a set of inputs $x \in \mathbb{R}^p$ and corresponding targets $y \in \mathbb{R}^d$. These serve as the training data on which the neural network is trained. We will adopt the terminology used by the TensorFlow framework [15] to help make the transition from mathematics to code easier.

6.1.1 Basic Mathematical Structure

A neural network is most generally defined as a non-linear function $f : \mathbb{R}^p \rightarrow \mathbb{R}^d$ built up as follows.

- A set of L layers. Consider the ℓ 'th layer. It consists of n_ℓ nodes all of which has a one-to-one correspondence to a real number. The conventional representation is with a real-valued vector $a^\ell \in \mathbb{R}^{n_\ell}$ called the *activation* of layer ℓ .
- For convenience, the layer with $\ell = 1$ is often called the *input layer* and the layer with $\ell = L$ is referred to as the *output layer*. The layers in between for $\ell = 2, \dots, L - 1$ are called the *hidden layers*. Although this distinction is merely conceptual and does not change the mathematics one bit, it provides useful categories for discussion later on.
- Each layer ℓ is supplied with a (possibly) non-linear function $\sigma_\ell : \mathbb{R}^{n_{\ell-1}} \rightarrow \mathbb{R}^{n_\ell}$. In other words, it defines a mapping $a^{\ell-1} \mapsto a^\ell$. The complete neural network function can thus be expressed as

$$f(x) = (\sigma_L \circ \sigma_{L-1} \circ \dots \circ \sigma_\ell \circ \dots \circ \sigma_2 \circ \sigma_1)(x). \quad (6.1)$$

- To each layer, we assign a *kernel* $W^\ell \in \mathbb{R}^{n_\ell \times n_{\ell-1}}$ and a *bias* $b^\ell \in \mathbb{R}^{n_\ell}$. Together, these parameters are called the *weights* of layer ℓ .
- The complete set of neural network parameters $(W, b) \equiv \{(W^\ell, b^\ell)\}_{\ell=1}^L$ are called the weights of the network. They serve as the *learnable* or *trainable* parameters of the model.
- Finally, we introduce the *logits* $z^\ell \in \mathbb{R}^{n_\ell}$ of layer ℓ .
- The permutation of number of layers, number of nodes per layer and activation functions are collectively called the *architecture* of the neural network.

The activation in layer ℓ is computed through the recursive equation:

$$a_j^\ell = \sigma_\ell \left(\sum_k W_{jk}^\ell a_k^{\ell-1} + b_j^\ell \right) \equiv \sigma_\ell(z_j^\ell), \quad \text{for } j = 1, 2, \dots, n_\ell. \quad (6.2)$$

A special case of eq. (6.2) applies to $\ell = 1$ where $a^0 = x \in \mathbb{R}^p$ is assumed.

6.1.2 Backpropagation

The standard approach to train a neural network is by minimization of some loss function by employing the backpropagation algorithm [16] to compute its gradient with respect to its trainable parameters recursively. The algorithm boils down to four equations. Consider \mathcal{L} as the loss function. The first of the four equations quantifies the error in the output layer.

$$\Delta_j^L = \frac{\partial \mathcal{L}}{\partial z_j^L}. \quad (6.3)$$

The second equation allows us to compute the error at layer ℓ given we know the error at layer $\ell + 1$,

$$\Delta_j^\ell = \left(\sum_k \Delta_k^{\ell+1} W_{kj}^{\ell+1} \right) \sigma'_\ell(z_j^\ell). \quad (6.4)$$

The final two equations relate these errors to the gradient of the loss function with respect to the model parameters. For the kernels, we have

$$\frac{\partial \mathcal{L}}{\partial W_{jk}^\ell} = \frac{\partial \mathcal{L}}{\partial z_j^\ell} \frac{\partial z_j^\ell}{\partial W_{jk}^\ell} = \Delta_j^\ell a_k^{\ell-1}. \quad (6.5)$$

For the biases, the gradients are

$$\frac{\partial \mathcal{L}}{\partial b_j^\ell} = \frac{\partial \mathcal{L}}{\partial z_j^\ell} \frac{\partial z_j^\ell}{\partial b_j^\ell} = \Delta_j^\ell. \quad (6.6)$$

With these four equations, we can fit the neural network using minimization techniques such as stochastic gradient descent or more complex methods such as ADAM (pages 13-19 in [3]). Although not the focus of this thesis, we might use these methods in conjunction with HMC to speed up convergence to the stationary distribution. Furthermore, the computation of gradients in combination with HMC or NUTS is achieved with the backpropagation algorithm.

We are now equipped to write down the backpropagation for a single datapoint. It's built up of a *forward pass* which takes an input x and applies the recursive eq. (6.2) which produces a model prediction $\hat{y} = a^L$. The second part of the algorithm is the *backward pass* which based on the prediction \hat{y} and the target y , computes the gradients of the loss function \mathcal{L} with respect to the model parameters. The forward pass of the neural network is summarized algorithm 6.1.

Algorithm 6.1 Backpropagation: Forward pass

```

procedure FORWARDPASS( $x$ )
   $a_j^0 = x_j$    for  $j = 1, \dots, p$                                  $\triangleright$  Initialize input
  for  $\ell = 1, 2, \dots, L$  do
    for  $j = 1, 2, \dots, n_\ell$  do
       $a_j^\ell \leftarrow \sigma_\ell \left( \sum_k W_{jk}^\ell a_k^{\ell-1} + b_j^\ell \right)$ 
    end for
  end for
end procedure

```

The backward pass of the algorithm is stated in algorithm 6.2.

Algorithm 6.2 Backpropagation: Backward pass

```

procedure BACKWARDPASS( $\mathcal{L}, x, y$ )
  for  $j = 1, 2, \dots, n_L$  do
     $\Delta_j^L \leftarrow \partial \mathcal{L} / \partial z_j^L$ 
     $\partial \mathcal{L} / \partial b_j^L \leftarrow \Delta_j^L$ 
     $\partial \mathcal{L} / \partial W_{jk}^L \leftarrow \Delta_j^L a_k^{L-1}$ 
  end for
  for  $\ell = L - 1, \dots, 1$  do
    for  $j = 1, \dots, n_\ell$  do
       $\Delta_j^\ell \leftarrow \left( \sum_k \Delta_k^{\ell+1} W_{kj}^{\ell+1} \right) \sigma'(z_j^\ell)$ 
       $\partial \mathcal{L} / \partial b_j^\ell \leftarrow \Delta_j^\ell$ 
       $\partial \mathcal{L} / \partial W_{jk}^\ell \leftarrow \Delta_j^\ell a_k^{\ell-1}$ 
    end for
  end for
end procedure

```

Note that for in all practical implementations in this thesis, we utilize *automatic differentiation* provided by TensorFlow to compute the gradients.

6.1.3 Regularization in Neural Networks

As discussed in chapter 2, models with a large number of parameters are prone to overfit training data and generalize poorly as a consequence. Thus one typically tack on an L^2 -regularization term to the loss \mathcal{L}_0 . Assuming that \mathcal{L}_0 is the RSS in eq. (2.2), the form of the full loss function for a neural network model becomes

$$\mathcal{L} = \frac{1}{2} \sum_i \left\| \hat{y}^{(i)} - y^{(i)} \right\|_2^2 + \frac{\lambda_W}{2} \sum_\ell \|W^\ell\|_2^2 + \frac{\lambda_b}{2} \sum_\ell \|b^\ell\|_2^2, \quad (6.7)$$

where λ_W and λ_b are regularization strengths for the kernels and biases respectively. The L^2 -norm $\|\cdot\|_2$ is the standard Euclidean norm in the case of a vector. For a matrix, we mean the following. Let $A \in \mathbb{R}^{m \times n}$. The matrix norm $\|\cdot\|_2$ is then given by *Fröbenius norm*

$$\|A\|_2 = \sqrt{\sum_{i=1}^m \sum_{j=1}^n |A_{ij}|^2}. \quad (6.8)$$

6.2 Activation Functions

There are many common activation functions with various strengths used in modern neural networks. We shall briefly mention a few ones for completeness.

6.2.1 Sigmoid and Tanh

The sigmoid activation function is given by

$$\sigma(x) = \frac{1}{1 + \exp(-x)}. \quad (6.9)$$

It was a very common choice in neural networks early, likely due to its simple derivative. It has a significant drawback, however. Looking at eq. (6.9), we can easily deduce that $\sigma(\pm\infty) = 0$, and since its derivative is of the form $\sigma'(x) = \sigma(1 - \sigma)$, the gradient computed during backpropagation vanishes if the input to the activation function as $|x| \rightarrow \infty$. This significantly hampers the progress during optimization. A popular alternative to the sigmoid function is the hyperbolic tangent given by

$$\tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1}. \quad (6.10)$$

This function is very similar to sigmoid in the sense that its derivative vanishes for inputs of large magnitude and so may suffer from the same issues as sigmoid does.

6.2.2 ReLU

To overcome the vanishing gradient problem, an activation function called the Rectifying Linear Unit (ReLU) became widely adopted, which is given by

$$\sigma(x) = x^+ = \max(0, x). \quad (6.11)$$

6.2.3 Swish

Recently, an activation function to replace ReLU was proposed in [17] known as *swish* or SiLU which was shown to outperform ReLU in deep neural networks on a number of challenging datasets. The activation function is given by

$$\sigma(x) = x \cdot \text{sigmoid}(x). \quad (6.12)$$

6.3 Bayesian learning of Neural Networks using Monte Carlo Samplers

So far, we have discussed neural networks as a model class whilst ignoring the issue of what it really means to do Bayesian learning of neural networks, in other words, what it means to *train* BNNs. We have left it somewhat unclear what this means and as it turns out, its meaning can be quite different depending on how Bayesian inference is performed. In this section we will clarify precisely what it means to train BNN using MCMC samplers such as HMC and NUTS. We shall then discuss practical aspects of the training which we shall put to practice in chapter 7.

6.3.1 What *is* Bayesian learning of Neural Networks?

The way Bayesian learning of neural networks manifest itself depends on the way in which we do Bayesian inference of the probabilistic model. We are concerned with inference of model parameters from the posterior using MCMC methods and will therefore obtain samples where each such sample consist of the weights of an entire neural network. More precisely, if we gather N samples with a chosen sampler, we will obtain N entire neural networks all sampled from the posterior to explain the observed data. Thus, what we mean by a *trained* BNN in this sense is that we have sampled a set of neural networks that collectively represent the BNN.

As we discussed at the end of chapter 2, we are mainly interested in the predictive distribution $p(y|x, W, b)$ of an output y given an input x . We can approximate this distribution by constructing an empirical distribution by feeding x through all N sampled neural networks to obtain N predicted targets \hat{y} using eq. (2.26). The second quantity of interest is expectations of target functions dependent on the model parameters. We can approximate any such expectation with an MCMC estimator as in eq. (3.4) using all N networks to evaluate the target function.

6.3.2 The Potential Energy Function of Neural Networks

We now turn to the Bayesian formulation of the neural network model for use with the samplers used in this thesis. Assume that we have picked an architecture for a neural network and wish to train it in the Bayesian sense. For both HMC and NUTS, we need only specify a potential energy function for our model. The samplers take care of the rest. Assume we are dealing with a dataset $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$ where all N points are independent and identically distributed. Equation (4.15) instructs us to specify a prior for the weights of the network, and a likelihood function that depends on the target and the model output, in order to fully specify the potential energy function. Common practice is to choose priors that are either Gaussian or Laplacian. We will operate with Gaussian priors, i.e

$$P(W^\ell) \propto \exp\left(-\frac{\lambda_W}{2}\|W^\ell\|_2^2\right) \quad \text{and} \quad P(b^\ell) \propto \exp\left(-\frac{\lambda_b}{2}\|b^\ell\|_2^2\right). \quad (6.13)$$

We will not worry too much about the choice of priors as the term in the potential energy function that corresponds to the likelihood will be much larger in practice. The Gaussian priors serve roughly the same purpose as L^2 -regularization does in classical ML.

The likelihood for regression from eq. (2.17) formulated in terms of a neural network $\hat{f}(x^{(i)}; W, b)$ is

$$p(D|\theta) = \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N \|y^{(i)} - \hat{f}(x^{(i)}; W, b)\|_2^2\right). \quad (6.14)$$

This is not the only valid choice for a likelihood function but it is the common choice since it can be identified with Euclidean L^2 -norm and has “neat” mathematical properties.

Combining the priors and the likelihood with eq. (2.24) yields the potential energy function

$$\mathcal{L} = \frac{1}{2\sigma^2} \sum_{i=1}^N \|y^{(i)} - f(x^{(i)}; W, b)\|_2^2 + \frac{\lambda_W}{2} \sum_{\ell=1}^L \|W^\ell\|_2^2 + \frac{\lambda_b}{2} \sum_{\ell=1}^L \|b^\ell\|_2^2, \quad (6.15)$$

up to a constant. As we discussed in chapter 4, the potential energy function also happens to be the typical loss function with L^2 -regularization used in the classical ML which is why we denote it as \mathcal{L} . At this point, we have set up all the machinery we need to train BNNs. Our next topic of discourse is the practice of doing so.

6.3.3 Practical Training of Bayesian Neural Networks

Training BNNs in practice requires us to specify a fairly large number of hyperparameters to obtain a set models. These are

1. **Neural network architecture.** We need to specify its number of layers, number of nodes and activation function per layer. Once the BNN is trained, we store this information along with the model for future usage. The stored weights themselves will encode how many layers and nodes the model has but the activation functions must be stored in addition.
2. **Number of results.** We must specify how many neural networks we want to sample and store. Because the weights must be stored in its entirety, we are forced to worry about the amount of disk space that is required to do so. For a fixed allocated disk space, we can obviously store a larger set of samples if the model is simple. As complexity increases, the number of samples we can store will necessarily decrease.
3. **Number of burn-in steps.** We must decide how long we want to run the MCMC chain before we start storing results. If disk space was no object, this step would be considered entirely optional as we could simply store every single sample and make a thorough analysis of the chain’s quality to determine when proper mixing is obtained. In practice, with TensorFlow’s framework, we can make a predetermined set of burn-in steps to avoid unnecessary RAM usage.

4. **Amount of thinning.** Since successive samples most likely will be correlated, we can specify how many samples we simply skip once we start gathering samples, i.e after the burn-in period. Again, we could ignore this and do this manually with the chain but doing so becomes a question of amount of available VRAM, RAM and disk space.
5. **Hyperparameters specific to the samplers.** The samplers themselves carry their own hyperparameters. In the case of HMC, we must specify a fixed number of Leapfrog steps L . If we use the NUTS sampler, we must specify the maximum tree depth. Moreover, we must determine how much of the computing resources we allocate to adapting the step size used in the Leapfrog integrator.
6. **Amount of pretraining.** An attempt to accelerate convergence of the MCMC chain can be achieved by pretraining the neural network using minimization methods with the backpropagation algorithm to bring the weights closer to a minima of the potential energy function (i.e the loss function used in classical ML). Then the point estimate obtained at the end of the training is used as a starting point for the MCMC chain.

6.3.4 Training Algorithm of Bayesian Neural Networks

In this section we shall turn our attention to an actual training algorithm for BNNs. Assume we pick a sampler S that represents either HMC or NUTS and a specified permutation of the hyperparameters discussed in the last section. In practice we can summarize a training algorithm as follows.

1. Initialize the weights of the model from the specified priors, i.e

$$W^\ell \sim p(W^\ell) \quad \text{and} \quad b^\ell \sim p(b^\ell) \quad \text{for} \quad \ell = 1, \dots, L. \quad (6.16)$$

2. Minimize the potential energy function \mathcal{L} with respect to the weights of the model using an optimizer of your choice to obtain a point estimate for use as the initial state of the Markov chain.
3. Initialize the Markov chain for a finite set of burn-in steps to achieve mixing using S . A proportion of the initial burn-in steps are used for step size adaptation, while the remaining are used for mixing.
4. Gather samples by applying S repeatedly, replacing the current weights of the model by the ones returned by S .

Chapter 7

Numerical Experiments

7.1 The Dataset

7.1.1 Data Generation

7.1.2 Data Scaling and Transformations

We shall briefly discuss how the training data is transformed before training. The targets in the dataset of NLO cross sections can span several orders of magnitude. For practical training of BNNs, this would require model parameters that also span several orders of magnitude. The result will usually be overflow and thus unsuccessful training of the models. Therefore, we have chosen to map the targets using the base-10 logarithm, i.e $y \mapsto \log_{10}(y)$. More generally, we could choose any base- a logarithm. A practical consideration here is that once the model is trained, any prediction it produces must be transformed back using the inverse mapping. As we increase the value of a , the precision the model can predict decreases. Thus a small error in log-space can result in a large error in what we may refer to as target space the larger the value of a is.

7.2 Performance Metrics

In this section, we will discuss the performance metrics used to benchmark and measure the performance of the models trained in this thesis. Due to the inherent probabilistic nature of the models trained, any output the model produces will be a distribution from which we can calculate a sample mean and variance.

7.2.1 Relative Error

The first and simplest form of performance metric we can use is the *relative error* which is defined as

$$\epsilon(x^*) = \frac{y_{\text{true}} - \hat{y}_{\text{mean}}(x^*)}{y_{\text{true}}}, \quad (7.1)$$

where y_{true} is the true target and $\hat{y}_{\text{mean}}(x^*)$ is the sample mean of the empirical predictive distribution of the model.

7.2.2 Standardized Residuals

A particularly useful way to represent how well a probabilistic model performs is to study its *standardized residual* which is given by

$$z(x^*) = \frac{y_{\text{true}} - \hat{y}_{\text{mean}}(x^*)}{\hat{\sigma}(x^*)}, \quad (7.2)$$

where $\hat{\sigma}(x^*)$ is the square-root of the sample variance. The mathematical representation of the targets y is that they can be decomposed as

$$y = f(x) + \delta, \quad (7.3)$$

for some true function $f(x)$ and a random noise $\delta \sim \mathcal{N}(0, 1)$, i.e it is distributed according to a standard Normal distribution. But in the case of data produced by **Prospino**, the noise is negligible which means that $y \approx f(x)$. The regression error obtained through the sample variance is therefore dominated by the predictive distribution computed by the model itself. A good model

7.3 Results

7.3.1 Benchmarks of Hyperparameters

7.3.1.1 Baseline Model

7.3.1.2 Pretraining

7.3.1.3 Burn-in length

7.3.1.4 Number of model parameters

7.3.2 Neutralino-Neutralino Cross Sections

Burn-in steps	Number of steps between	Kernel	Number of results	Pretraining epochs	Pretraining batch size
2500	10	NUTS	1000	1000	32

Table 7.1: The table shows the training configuration used to sample the models listed in table 7.2.

Model number	Architecture	Number of parameters
1	5-50-1	351
2	5-50-50-1	2901
3	5-50-50-50-1	5451
4	5-50-50-50-50-1	8001
5	5-50-50-50-50-50-1	10551

Table 7.2: The table shows the models used in this section. For each model, 1000 sampled networks are used to generate each result shown in this section. The architecture describe number of nodes per layer. For each hidden layer, the same activation function is used. The final layer uses an identity function.




Figure 7.1: The figure shows histograms of the standardized residuals computed for different BNNs. The normal distribution is drawn as dotted line.

Chapter 8

Discussion

Conclusion

Conclusion here.

Appendices

Appendix A

A.1 Appendix 1 title

Some appendix stuff.

Bibliography

- [1] D. P. Kingma and J. Ba, *Adam: A method for stochastic optimization*, 2014. 10.48550/ARXIV.1412.6980.
- [2] J. L. Devore and K. N. Berk, *Modern Mathematical Statistics with Applications*, p. 80. Springer, 2018.
- [3] P. Mehta, M. Bukov, C.-H. Wang, A. G. Day, C. Richardson, C. K. Fisher et al., *A high-bias, low-variance introduction to machine learning for physicists*, *Physics Reports* **810** (may, 2019) 1–124.
- [4] M. Betancourt, *A conceptual introduction to hamiltonian monte carlo*, 2017. 10.48550/ARXIV.1701.02434.
- [5] G. O. Roberts and J. S. Rosenthal, *General state space markov chains and MCMC algorithms*, *Probability Surveys* **1** (jan, 2004) .
- [6] A. Gelman and D. B. Rubin, *Inference from Iterative Simulation Using Multiple Sequences*, *Statistical Science* **7** (1992) 457 – 472.
- [7] S. Brooks, A. Gelman, G. L. Jonas and X.-L. Meng, eds., *Handbook of Markov Chain Monte Carlo*, ch. 6. Springer, 2018.
- [8] N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller and E. Teller, *Equation of State Calculations by Fast Computing Machines*, *The Journal of Chemical Physics* **21** (1953) 1087–1092, [<https://doi.org/10.1063/1.1699114>].
- [9] W. K. Hastings, *Monte Carlo sampling methods using Markov chains and their applications*, *Biometrika* **57** (04, 1970) 97–109, [<https://academic.oup.com/biomet/article-pdf/57/1/97/23940249/57-1-97.pdf>].
- [10] J. S. Helbert Goldstein, Charles Poole, *Classical Mechanics*, 3rd ed., ch. 2,8. Addison Wesley, 2000.
- [11] C. M. Bishop, *Pattern Recognition and Machine Learning*, ch. 11, p. 551. Springer New York, 2006.
- [12] M. D. Hoffman and A. Gelman, *The no-u-turn sampler: Adaptively setting path lengths in hamiltonian monte carlo*, 2011.
- [13] Y. Nesterov, *Primal-dual subgradient methods for convex problems*, *Mathematical Programming* **120** (Aug, 2009) 221–259.
- [14] S. Brooks, A. Gelman, G. Jones and X.-L. Meng, eds., *Handbook of Markov Chain Monte Carlo*. Chapman and Hall/CRC, may, 2011, [10.1201/b10905](https://doi.org/10.1201/b10905).
- [15] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro et al., *TensorFlow: Large-scale machine learning on heterogeneous systems*, 2015.
- [16] D. E. Rumelhart, G. E. Hinton and R. J. Williams, *Learning representations by back-propagating errors*, *Nature* **323** (1986) 533–536.
- [17] P. Ramachandran, B. Zoph and Q. V. Le, *Searching for activation functions*, *CoRR* **abs/1710.05941** (2017) , [[1710.05941](https://arxiv.org/abs/1710.05941)].