Bayesian neural network estimation of next-to-leading-order cross sections

by

René Alexander Ask

THESIS

for the degree of

MASTER OF SCIENCE

Faculty of Mathematics and Natural Sciences
University of Oslo

Autumn 2021

Bayesian neural network estimation of next-to-leading-order cross sections

René Alexander Ask

Abstract

This is my abstract.

ii

Acknowledgments

Acknowledgments yo

Contents

Introduction

Motivation, context and problem.

Outline of the Thesis

Give outline of thesis

Chapter 1

Hamiltonian Monte Carlo

In this chapter, we will study the details of the Hamiltonian Monte Carlo (HMC) method. It is a Markov Chain Monte Carlo (MCMC) sampling technique that merges Gibbs sampling and a modified version of Metropolis sampling with Hamiltonian dynamics. It avoids the random walk behaviour of the Metropolis algorithm and generates successive samples with smaller correlation. In this chapter, the first thing we will discuss is Markov chains. From that, follows an outline of Gibbs- and Metropolis sampling. After that, we will deal with Hamiltonian dynamics before we bring it all together to form the Hamiltonian Monte Carlo algorithm for sampling. This forms the basis for the practical Bayesian learning of the neural networks studied in this thesis.

1.1 Markov Chain Monte Carlo

Monte Carlo Markov chain (MCMC) is a scheme to sample points θ proportional to a distribution $\pi(\theta)$. It generates a new point θ_i given a point θ_{i-1}. By Markov chain, we mean a sequence of points $\theta_1, \theta_2, \ldots$ that are possibly dependent, but are occuring in the sequence in proportion to $\pi(\theta)$. Note that $\pi(\theta)$ here is not an exact probability distribution because it need not be normalized to unity. However, suppose $P(\theta)$ is the underlying probability distribution, then $\pi(\theta) \propto P(\theta)$. Typically, in Bayesian applications, we have a prior $P(\theta)$ and a likelihood $P(D|\theta)$. In this case $\pi(\theta) = P(D|\theta)P(\theta)$ and $\pi(\theta) \propto P(\theta|D)$, that is, it's proportional to the posterior distribution.

A few important properties of the Markov chain, originally introduced by Metropolis et. al and built upon by Hastings [1]:

1. **Ergodicity**. Each point θ_i is chosen from a distribution that only depends on the previous point in the sequence, θ_{i-1}. For this, we introduce a transition probability that is $T(\theta_i|\theta_{i-1})$. This ensures that any point θ can eventually be reached given a long enough sequence of samples [2].

2. **Detailed balance**. The transition probability is chosen to obey
$$\pi(\theta)T(\theta'|\theta) = \pi(\theta')T(\theta|\theta'),$$
which ensures that the Markov chain is ergodic. Mathematically, we can express this condition as
$$\pi(\theta') = \int \pi(\theta)T(\theta'|\theta)\mathrm{d}\theta.$$

3. We allow the transition $\theta \to \theta$, hence the transition probability $T(\theta|\theta)$ may be non-zero.

4. The transition probablities integrate to unity, thus
$$\int T(\theta'|\theta)\mathrm{d}\theta = 1.$$

5. Finally, the transition probabilities are required to be time–independent.

1.2   Gibbs sampling

Gibbs sampling [3] is a sampling technique used to generate a Markov chain sequence from an underlying multivariate distribution $P(\gamma)$ for a multi-dimensional parameter $\gamma = (\gamma_1, \ldots \gamma_d) \in \mathbb{R}^d$, for $d > 1$. Suppose $\gamma^{(t)}$ is the parameters at iteration t. Then the parameters $\gamma^{(t+1)}$ at iteration $t+1$ are generated from $\gamma^{(t)}$ by the following procedure.

Algorithm 1 Gibbs sampling

---

procedure GIBBS($\gamma^{(t)}$)

Sample $\gamma_1^{(t+1)} \sim P(\gamma_1 | \gamma_2^{(t)}, ..., \gamma_d^{(t)})$

Sample $\gamma_2^{(t+1)} \sim P(\gamma_2 | \gamma_1^{(t+1)}, ..., \gamma_d^{(t)})$

$\vdots \qquad \vdots \qquad \vdots \qquad \vdots$

Sample $\gamma_d^{(t+1)} \sim P(\gamma_d | \gamma_1^{(t+1)}, ..., \gamma_{d-1}^{(t+1)})$

end procedure

Thus each new sample $\gamma_i^{(t+1)}$ is only dependent on the prior state of the other parameters through

$$\gamma_i^{(t+1)} \sim P(\gamma_i | \gamma_1^{(t+1)}, \ldots \gamma_{i-1}^{(t+1)}, \gamma_{i+1}^{(t)}, \ldots, \gamma_d^{(t)}), \tag{1.1}$$

which by definition makes it a Markov chain.

1.3   The Metropolis-Hastings algorithm

The Metropolis algorithm [1] is a sampling algorithm based on random walks in parameter space used in MCMC chains to generate a new point θ' given a point θ. Albeit efficient for some problems, it's not a suitable sampling technique in the context of neural networks. However, a rudimentary understanding of the algorithm will be useful before we embark upon the HMC sampling algorithm.

The transition probability in the Metropolis algorithm is chosen to be

$$T(\theta'|\theta) = q(\theta'|\theta)A(\theta, \theta'), \tag{1.2}$$

where $q(\theta'|\theta)$ is called the proposal distribution and $A(\theta, \theta')$ is the acceptance probability given by

$$A(\theta, \theta') = \min\left(1, \frac{\pi(\theta')q(\theta|\theta')}{\pi(\theta)q(\theta'|\theta)}\right). \tag{1.3}$$

The point θ' is accepted with probability $A(\theta, \theta')$.

Algorithm 2 Metropolis-Hastings

---

procedure METROPOLIS-HASTINGS(θ)

Sample $\theta' \sim q(\theta'|\theta)$

$p \leftarrow \min\left(1, \frac{\pi(\theta')q(\theta|\theta')}{\pi(\theta)q(\theta'|\theta)}\right)$

Sample u uniformly on $(0, 1)$.

if $p \geq u$ **then**
    $\theta \leftarrow \theta'$ ▷ Accept transition
else
    $\theta \leftarrow \theta$ ▷ Reject transition
end if

end procedure

1.4 Hamiltonian dynamics

Hamiltonian dynamics [4] plays a central part in the HMC algorithm. For completeness, we will first survey Lagrangian mechanics from which we derive the Hamiltonian. The Hamiltonian then lays the foundation for the Hamiltonian dynamics.

1.4.1 Lagrangian Mechanics

Assume a set of *generalized coordinates* $q = (q_1, ..., q_n)$. Generally, the Lagrangian can be written as

$$L(q, \dot{q}, t) = K(q, \dot{q}, t) - V(q, \dot{q}, t), \tag{1.4}$$

where K is the kinetic energy and V is the potential energy of the system. We shall restrict the treatment to the case where there is no explicit dependence on time t. The solutions $q(t)$ can be found by solving the *Euler–Lagrange* equations given by

$$\frac{\mathrm{d}}{\mathrm{d}t}\frac{\partial L}{\partial \dot{q}_i} - \frac{\partial L}{\partial q_i} = 0. \tag{1.5}$$

1.4.2 Hamiltonian Mechanics

The Hamiltonian is constructed by the Legendre transformation,

$$H(q, p, t) = \sum_i p_i \dot{q}_i(q, p) - L(q, \dot{q}(q, p), t), \tag{1.6}$$

where

$$p_i = \frac{\partial L}{\partial \dot{q}_i}. \tag{1.7}$$

The equations of motion, known as *Hamilton's* equations, are given by

$$\frac{\mathrm{d}q_i}{\mathrm{d}t} = \frac{\partial H}{\partial p_i}, \qquad \frac{\mathrm{d}p_i}{\mathrm{d}t} = -\frac{\partial H}{\partial q_i}. \tag{1.8}$$

For the purpose of utilizing this framework in the context of HMC, it's assumed that the form of the Lagrangian is

$$L(q, \dot{q}) = K(\dot{q}) - V(q), \tag{1.9}$$

where

$$K(\dot{q}) = \sum_i \frac{1}{2} m_i \dot{q}_i^2. \tag{1.10}$$

The generalized momentum of coordinate q_i is

$$p_i = \frac{\partial K}{\partial \dot{q}_i} = m\dot{q}_i, \tag{1.11}$$

from which it follows that the Legendre transformed kinetic energy can be written as

$$K(p) = \sum_i \frac{p_i^2}{2m_i}. \tag{1.12}$$

Finally, we can write down the Hamiltonian as

$$H(q,p) = K(p) + V(q) = \sum_i \frac{p_i^2}{2m_i} + V(q). \tag{1.13}$$

1.4.3 Leapfrog integration

To run one step of HMC, we need to compute the time evolution of a Hamiltonian system of the form discussed in the former section, where the neural network parameters will play the role as the generalized coordinates q. The common choice of algorithm to integrate the equations of motion in eq. (1.8) is *Leapfrog* integration [5]. This integrator is *symplectic*, which means it conserves local volumes in phase space. This effectively translates to an approximately conserved value of $H(q,p)$ throughout a simulation, with slight oscillations about a mean value.

Assume we approximate the true coordinates and momenta by $(\hat{q}, \hat{p})$. A single Leapfrog integration step can then be written as in algorithm 3. Here h represents the stepsize used in the algorithm.

Algorithm 3 Leapfrog integration (single step)

---

procedure LEAPFROG(q, p, λ)
    1. $\hat{p}_i(t + h/2) = \hat{p}_i(t) - \lambda\frac{h}{2} \, \partial V(\hat{q}(t))/\partial q_i$

    2. $\hat{q}_i(t + h) = \hat{q}_i(t) + \lambda\frac{h}{m_i}\hat{p}_i(t + h/2)$

    3. $\hat{p}_i(t + h) = \hat{p}_i(t + h/2) - \lambda\frac{h}{2} \, \partial V(\hat{q}(t + h))/\partial q_i$
end procedure

1.5 Hamiltonian Monte Carlo

Hamiltonian Monte Carlo is largely developed and expanded upon by Radford Neal [6]. The probability distribution to sample from is expressed in terms the Canonical distribution

$$P(q) \propto \exp(-V(q)), \tag{1.14}$$

where q represents the parameters of the model. $V(q)$ can then always be expressed as

$$V(q) = -\log Z - \log P(q) \tag{1.15}$$

for some Z. To utilize the framework of Hamiltonian dynamics explained in section 1.4, we introduce momentum variables p_i such that we can express the total Hamiltonian as

$$H(q,p) = K(p) + V(q), \tag{1.16}$$

with a corresponding canonical distribution over phase-space

$$P(q, p) \propto \exp\left(-H(q, p)\right). \tag{1.17}$$

The algorithm is summarized in

Algorithm 4 Hamiltonian Monte Carlo

---

  procedure HMC(L, q, p)
    Sample $u \sim U(0, 1)$.
    $\lambda = 1$ if $u \geq 1/2$ else $\lambda = -1$
    $(q^*, p^*) \leftarrow (q, p)$ $\triangleright$ Start from initial state.
    $p^* \leftarrow \text{GIBBS}(p^*)$
    for $l = 1, ..., L$ **do** $\triangleright$ L Leapfrog steps.
        $(q^*, p^*) \leftarrow \text{LEAPFROG}(q^*, p^*, \lambda)$
    end for
    $P = \min\left(1, \exp\left[-\left(H(q^*, p^*) - H(q, p)\right)\right]\right)$
    Sample $u \sim U(0, 1)$ $\triangleright$ Uniform distribution on $(0, 1)$.
    if $P \geq u$ **then**
        $(q, p) \leftarrow (q^*, p^*)$ $\triangleright$ Accept proposed state.
    else
        $(q, p) \leftarrow (q, p)$ $\triangleright$ Reject proposed state.
    end if
  end procedure

Chapter 2

Bayesian Learning for Neural Networks

2.1 Neural Networks

In this chapter, we will introduce the mathematical formalism underpinning neural networks. For convenience, we will adopt the terminology used by Tensorflow[7] to help make the transition from the mathematics to their machine learning framework easier. We will stay general and assume a set of inputs $x \in \mathbb{R}^p$ and corresponding targets $y \in \mathbb{R}^d$. These serve as the training data on which the neural network is trained.

2.1.1 Basic mathematical structure

A neural network is most generally defined as a non-linear function $f : \mathbb{R}^p \to \mathbb{R}^d$. This non-linear function is built-up as follows:

- A set of L layers. Consider the ℓ'th layer. It consists of n_ℓ nodes all of which has a one-to-one correspondence to a real number. The conventional representation is through a real-valued vector $a^\ell \in \mathbb{R}^{n_\ell}$, where a^ℓ is colloquially called the *activation* of layer ℓ.

- For convenience, the layer with $\ell = 1$ is often called the *input layer* and the layer with $\ell = L$ is called the *output layer*, and the layers in between for $\ell = 2, ..., L - 1$ are called the *hidden layers*. Although this distinction is merely conceptual and does not change the mathematics one bit, it provides useful categories for discussion later on.

- Each layer ℓ is supplied with a (possibly) non-linear function $\sigma_\ell : \mathbb{R}^{n_{\ell-1}} \to \mathbb{R}^{n_\ell}$. In other words, it defines a mapping $a^{\ell-1} \mapsto a^\ell$. The complete neural network function can thus be expressed as

$$f(x) = (\sigma_L \circ \sigma_{L-1} \circ \cdots \circ \sigma_\ell \circ \cdots \circ \sigma_2 \circ \sigma_1)(x). \tag{2.1}$$

- To each layer, we assign a *kernel* $W^\ell \in \mathbb{R}^{n_\ell \times n_{\ell-1}}$ and a *bias* $b^\ell \in \mathbb{R}^{n_\ell}$. Together, these parameters are called the *weights* of a layer.

- The complete set of neural network parameters $(W, b) = \{(W^\ell, b^\ell)\}_{\ell=1}^L$ are called the weights of the network. They serve as the *learnable* or *trainable* parameters of the model.

- Finally, we introduce the *logits* $z^\ell \in \mathbb{R}^{n_\ell}$ of layer ℓ.

- The permutation of chosen number of layers, number of nodes per layer and activation functions are collectively called the *architecture* of the neural network.

The activation in layer ℓ is computed through the recursive equation:

$$a_j^\ell = \sigma_\ell\left(\sum_k W_{jk}^\ell a_k^{\ell-1} + b_j^\ell\right) \equiv \sigma_\ell(z_j^\ell), \quad \text{for} \quad j = 1, 2, ..., n_\ell. \tag{2.2}$$

A special case of eq. (2.2) applies to $\ell = 1$ where $a^0 = x \in \mathbb{R}^p$ is assumed.

2.1.2 Backpropagation

The standard approach to train a neural network is by minimization of some loss function by employing the *backpropagation* algorithm[8]. The algorithm boils down to four equations defining a recursive algorithm that approximates the gradient with respect to the parameters of the model. Consider E as the loss function, quanitifying the error between the target and the model output. The first of the four equations quantifes the error in the output layer.

$$\Delta_j^L = \frac{\partial E}{\partial z_j^L}. \tag{2.3}$$

The second equation allows us to compute the error at layer ℓ given we know the error at layer $\ell + 1$,

$$\Delta_j^\ell = \left(\sum_k \Delta_k^{\ell+1} W_{kj}^{\ell+1}\right) \sigma_\ell'(z_j^\ell). \tag{2.4}$$

The final two equations relate these errors to the gradient of the loss function with respect to the model parameters. For the kernels, we have

$$\frac{\partial E}{\partial W_{jk}^\ell} = \frac{\partial E}{\partial z_j^\ell} \frac{\partial z_j^\ell}{\partial W_{jk}^\ell} = \Delta_j^\ell a_k^{\ell-1}. \tag{2.5}$$

For the biases, the gradients are

$$\frac{\partial E}{\partial b_j^\ell} = \frac{\partial E}{\partial z_j^\ell} \frac{\partial z_j^\ell}{\partial b_j^\ell} = \Delta_j^\ell. \tag{2.6}$$

With these four equations, we can fit the neural network using minimization techniques such as stochastic gradient descent or more complex methods such as ADAM (pages 13-19 in [9]). Although not the focus of this thesis, we might use these methods in conjunction with HMC to speed up convergence to the stationary distribution. Furthermore, the computation of gradients in combination with HMC can be performed with the backpropagation algorithm.

2.1.3 Loss function for regression

In this thesis, we are concerned with regression tasks. The activation function of the final layer σ_L is then just the identity function. The typical loss function chosen to solve regression tasks is the L_2-norm, which for a single output can be written as

$$E(y, \hat{y}) = \frac{1}{2}\|y - \hat{y}\|_2^2, \tag{2.7}$$

where $\hat{y}$ denotes the model output and y the ground-truth. Now, the model output in this case is $\hat{y}_j = a_j^L = z_j^L$. Therefore,

$$\Delta_j^L = \frac{\partial E}{\partial z_j^L} = a_j^L - y_j. \tag{2.8}$$

We are now equipped to write down the backpropagation for a single datapoint. It's built up of a textitforward pass which takes an input x and applies the recursive eq. (2.2) which produces a model prediction $\hat{y} = a^L$. The second part of the algorithm is the *backward pass* which based on the prediction $\hat{y}$ and the target y, computes the gradients of the loss function with respect to the model parameters. The forward pass of the neural network is found in algorithm 5.

Algorithm 5 Backpropagation: Forward pass

---

procedure FORWARDPASS(x)
$\quad a_j^0 = x_j \quad$ for $\quad j = 1, \ldots, p \qquad\qquad\qquad\qquad\qquad\qquad\qquad$ ▷ Initialize input
$\quad$**for** $\ell = 1, 2, .., L$ **do**
$\quad\quad$**for** $j = 1, 2, .., n_\ell$ **do**
$\quad\quad\quad a_j^\ell \leftarrow \sigma_\ell \left(\sum_k W_{jk}^\ell a_k^{\ell-1} + b_j^\ell \right)$
$\quad\quad$**end for**
$\quad$**end for**
end procedure

The backward pass of the algorithm is stated in algorithm 6.

Algorithm 6 Backpropagation: Backward pass

---

procedure BACKWARDPASS(y)
$\quad$**for** $j = 1, 2, \ldots, n_L$ **do**
$\quad\quad \Delta_j^L \leftarrow a_j^L - y_j$
$\quad\quad \partial E / \partial b_j^L \leftarrow \Delta_j^L$
$\quad\quad \partial E \big/ \partial W_{jk}^L \leftarrow \Delta_j^L a_k^{L-1}$
$\quad$**end for**
$\quad$**for** $\ell = L - 1, \ldots, 1$ **do**
$\quad\quad$**for** $j = 1, \ldots, n$ **do**
$\quad\quad\quad \Delta_j^\ell \leftarrow \left(\sum_k \Delta_k^{\ell+1} W_{kj}^{\ell+1} \right) \sigma'(z_j^\ell)$
$\quad\quad\quad \partial E / \partial b_j^\ell \leftarrow \Delta_j^\ell$
$\quad\quad\quad \partial E \big/ \partial W_{jk}^\ell \leftarrow \Delta_j^\ell a_k^{\ell-1}$
$\quad\quad\quad$ update b_j^ℓ and W_{jk}^ℓ.
$\quad\quad$**end for**
$\quad$**end for**
end procedure

2.1.4 Regularization

Neural networks often end up with a large number of parameters, which makes them prone to *overfit* training data. This means that the trained parameters of the model is adjusted to capture trends in the training data which may not represent the underlying process the model tries to learn. The consequence is that it *generalizes* poorly to new unseen data, i.e its predictions are poor. A typical strategy to avoid this problem, is to introduce some form of *regularization*. A common choice is L^2-*regularization*, which for a neural network tacks on two additional sums to the loss function as follows:

$$E = \frac{1}{2} \sum_i \left\| \hat{y}^{(i)} - y^{(i)} \right\|_2^2 + \frac{\lambda_W}{2} \sum_\ell \left\| W^\ell \right\|_2^2 + \frac{\lambda_b}{2} \sum_\ell \left\| b^\ell \right\|_2^2, \tag{2.9}$$

where λ_W and λ_b are regularization strengths for the kernels and biases respectively. The L^2-norm $\|\cdot\|_2$ is the standard Euclidean norm in the case of a vector. For a matrix, we mean the following. Consider a matrix $A \in \mathbb{R}^{m \times n}$. The matrix norm $\|\cdot\|_2$ is then given by *Fröbenius norm*

$$\|A\|_2 = \sqrt{\sum_{i=1}^m \sum_{j=1}^n |A_{ij}|^2}. \tag{2.10}$$

L^2-regularization is sometimes called L^2-penalty because it penalizes assignment of large values to the model parameters. Its effect is thus shrinkage of the parameter space where accessible minima may reside.

2.2   Interpretations of probability

2.3   Bayesian inference

The foundation for Bayesian inference is Bayes' theorem, which can be formulated as

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}. \tag{2.11}$$

where D is observed data and θ are the model parameters. Some useful terminology is in order. $P(\theta)$ is called the *prior* distribution and embodies our prior knowledge of θ before any new observations are considered. $P(D|\theta)$ is called the *likelihood* function and provides information about θ learned from observing the data D. The *posterior* distribution $P(\theta|D)$ models our belief about θ after the data D is observed. More succinctly, we can write Bayes' theorem as

$$P(\theta|D) \propto P(\theta|D)P(\theta), \tag{2.12}$$

because its rarely of interest, or tractable, to compute $P(D)$, known as the *evidence*.

The objective of Bayesian inference is to compute the *predictive* distribution for an unseen datapoint y^*, which can be expressed as the integral over all parameters of the posterior distribution weighted by the likelihood. Given a dataset of observations $D = \{y^{(1)}, \ldots y^{(n)}\}$, this implies

$$P(y^*|D) = \int P(y^*, \theta|D)\mathrm{d}\theta = \int P(y^*|\theta, D)P(\theta|D)\mathrm{d}\theta = \int P(y^*|\theta)P(\theta|D)\mathrm{d}\theta, \tag{2.13}$$

which loosely describes the probability of observing y^* given the observations in D.

2.4   Bayesian framework for neural networks

We can specialize the equations used in Bayesian inference for neural networks in the context of regression. The predictive distribution $P(y|x, \theta)$ seeks to model a function $f : \mathbb{R}^p \to \mathbb{R}^d$ that for a given $x \in \mathbb{R}^p$ produces an output $y \in \mathbb{R}^d$. In the infinite data limit, the distribution should be a Dirac delta function. For finite datasets, however, we instead seek a distribution of outputs given the input features.

Consider a set of observations $D = \{(x^{(1)}, y^{(1)}), ..., (x^{(n)}, y^{(n)})\}$, where $x^{(i)} \in \mathbb{R}^p$ are the input features and $y^{(i)} \in \mathbb{R}^d$ are the targets. The equation for the predictive distribution of y^* given an input x^* changes to

$$P(y^*|x^*, D) = \int P(y^*|x^*, \theta)P(\theta|D)\mathrm{d}\theta. \tag{2.14}$$

Assuming that the the observations in D are drawn independently, the likelihood function can be expressed as

$$P(D|\theta) = \prod_{i=1}^{n} P(y^{(i)}|x^{(i)}, \theta). \tag{2.15}$$

In the context of regression, the likelihood for a given observation (x, y) is commonly chosen to be

$$P(y|x, \theta) \propto \exp\left(-\frac{\lambda}{2}\|y - f(x; \theta)\|_2^2\right). \tag{2.16}$$

where $f(x; \theta)$ is the output of the neural network and λ is a hyperparameter typically chosen to be identical for all inputs (x, y) during training. The likelihood found eq. (2.16) is equivalent to using L^2-norm as a loss function with regularization strength λ when framed as a minimization problem, which we will see shortly.

In practice, however, we instead sample a set of network parameters $\{\theta_1, ..., \theta_n\}$ from the posterior distribution

$$P(\theta|D) \propto P(D|\theta)P(\theta), \tag{2.17}$$

from which we can produce a set of predictions $\{\hat{y}_1, \ldots \hat{y}_n\}$ from the neural network model

$$\hat{y}_i = f(x, \theta_i). \tag{2.18}$$

Given this set of predictions, we can compute the sample mean

$$\hat{y}_{\text{MLE}} = \frac{1}{n} \sum_i f(x, \theta_i), \tag{2.19}$$

which is an approximation to the *maximum likelihood estimate* (MLE). Furthermore, an estimate of the error is provided by the sample *covariance*

$$\text{Cov}(\hat{y}) = \frac{1}{n-1} \sum_i (\hat{y}^{(i)} - \hat{y}_{\text{MLE}})(\hat{y}^{(i)} - \hat{y}_{\text{MLE}})^T. \tag{2.20}$$

The diagonal terms of eq. (2.20) yields the sample *variance* of the components of $\{\hat{y}^{(i)}\}_{i=1}^n$.

2.5 Sources of uncertainty

The main motivation utilizing Bayesian inference to neural networks is to be able to quantify uncertainty. Uncertainty can be further divided into two categories:

- **Epistemic uncertainty**: the systematic uncertainty. It practice, this uncertainty can be quantified by the posterior of the model $P(\theta|D)$. Thus it quantifies the uncertainty in the model itself.

- **Aleatoric uncertainty**: the statistical uncertainty, related to the conditional distribution $P(y|x, \theta)$. This uncertainty quantifies the error due to random chance, and measures the random noise in the data.

In principle, in the infinite data limit, both errors goes to zero as there's only a single parametrization θ that explains all data D, and there's only a single output y for a given x and θ. However, neural networks as so-called *over-parametrized* models in the sense that they have many equivalent parametrizations. These occur due to underlying symmetries of the model, two of which are *weight-space symmetry* and *scaling symmetry*. The former means that we can permute parameters in two adjacent and produce the same output as before the permutation. The latter case occur when the network uses non-linear activation that obey $\sigma(\alpha x) = \alpha \sigma(x)$. This relation implies that we can rescale the parameters of two adjacent layers by a factor of α and $1/\alpha$ respectively and produce the same output. Thus a parametrization that explains a set of observations does not exist for sufficiently complex models.

2.6 Bayesian learning using HMC

To learn from the data D using HMC, we need to define a potential energy function $V(q)$ and a kinetic energy function $K(p)$. To this end, we need to specify

1. A prior for the network parameters, i.e the kernels W_{ij}^ℓ and the biases b_j^ℓ. A common choice is

$$P(W^\ell) \propto \exp\left(-\frac{\lambda_W}{2} \|W^\ell\|_2^2\right), \qquad P(b^\ell) \propto \exp\left(-\frac{\lambda_b}{2} \|b^\ell\|_2^2\right). \tag{2.21}$$

2. A distribution of the generalized momenta of the model. The kinetic energy in eq. (1.12) automatically imposes that

$$P(p) \propto \prod_i \exp\left(-\frac{p_i^2}{2m_i}\right), \tag{2.22}$$

since $K(p) = -\log P(p)$.

3. A likelihood function given an input (x, y),

from which we can define the potential energy as follows:

$$V(W, b) = -\log Z - \sum_{\ell} \log P(W^{\ell}) - \sum_{\ell} \log P(b^{\ell}) - \sum_{i} \log P(y^{(i)} | x^{(i)}, W, b), \tag{2.23}$$

where W denotes all the kernels and b denotes all the biases of the model. Here Z denotes the normalization constant of the distributino $P \propto \exp(-V(W, b))$, but is of no importance for the sampling procedure. Inserting the terms from eq. (2.21) and eq. (2.16), we get the following expression for the potential energy.

$$V(W, b) = -\log Z + \frac{\lambda_W}{2} \sum_{\ell} \left\| W^{\ell} \right\|_2^2 + \frac{\lambda_b}{2} \sum_{\ell} \left\| b^{\ell} \right\|_2^2 + \frac{\lambda}{2} \sum_{i} \left\| y^{(i)} - f(x^{(i)}; W, b) \right\|_2^2. \tag{2.24}$$

The terms in eq. (2.24) play the same role as in typical machine learning applications. Note that Z here is some appropriate constant that doesn't really matter for the sampling process,

With the ingredients introduced hitherto, we can proceed to sample from a network with an arbitrary network architecture.

Conclusion

Conclusion here.

Appendices

Appendix A

A.1   Appendix 1 title

Some appendix stuff.

Bibliography

[1] N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller and E. Teller, *Equation of State Calculations by Fast Computing Machines*, *The Journal of Chemical Physics* **21** (1953) 1087–1092, [https://doi.org/10.1063/1.1699114].

[2] W. H. P. et. al, *Numerical recipes in C : the art of scientific computing*, pp. 824–826. Second edition. Cambridge [Cambridgeshire] ; New York : Cambridge University Press, 1992., 1992.

[3] S. Geman and D. Geman, *Stochastic relaxation, Gibbs Distributions, and the Bayesian Restoration of Images*, *IEEE Transactions on Pattern Analysis and Machine Intelligence* **PAMI-6** (1984) 721–741.

[4] J. S. Helbert Goldstein, Charles Poole, *Classical Mechanics, 3rd ed.*, ch. 2,8. Addison Wesley, 2000.

[5] C. M. Bishop, *Pattern Recognition and Machine Learning*, ch. 11, p. 551. Springer New York, 2006.

[6] R. Neal, *Handbook of Markov Chain Monte Carlo*, ch. 5. Chapman and Hall/CRC, May, 2011. http://dx.doi.org/10.1201/b10905. 10.1201/b10905.

[7] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro et al., *TensorFlow: Large-scale machine learning on heterogeneous systems*, 2015.

[8] D. E. Rumelhart, G. E. Hinton and R. J. Williams, *Learning representations by back-propagating errors*, *Nature* **323** (1986) 533–536.

[9] P. Mehta, M. Bukov, C.-H. Wang, A. G. Day, C. Richardson, C. K. Fisher et al., *A high-bias, low-variance introduction to machine learning for physicists*, *Physics Reports* **810** (May, 2019) .